

Perceive, Plan, Act, Self-Correct: An Architectural Framework for Goal-Directed Agentic AI Systems

[Removed for double-blind review]

Abstract

Large language model (LLM) agents that autonomously perceive environments, plan multi-step strategies, execute actions through external tools, and self-correct from feedback represent a paradigm shift from prompt-response to goal-directed AI. Yet the field lacks a unified architectural vocabulary, hindering principled comparison. Here we propose PERCEIVE-PLAN-ACT-SELF-CORRECT (PPAS), a four-phase canonical loop grounded in classical agent theory and extended for LLM-era systems. We decompose agentic architectures into an 8-layer technology stack, mapping 15+ frameworks to a common reference. Three empirical studies—a benchmark meta-analysis across SWE-Bench Verified, WebArena, and GAIA (12 agents), a controlled design-pattern comparison, and an inter-agent protocol evaluation (MCP, A2A, AG-UI)—validate the framework. Reflective planning with human-in-the-loop approval achieves the highest task-completion rates (72.5% SWE-Bench) while reducing irreversible-action failures by 73%. We present a failure-mode taxonomy and EU AI Act governance mapping for responsible deployment.

Keywords: Agentic AI, Large language models, AI architectures, Multi-agent systems, Agent frameworks, Human-in-the-loop

1. Introduction

The rapid evolution of large language models (LLMs) has given rise to *agentic AI* systems that autonomously perceive their environment, formulate multi-step plans, execute actions via external tools, and self-correct from feedback

26 [79, 62, 70]. Unlike static generative models, these systems maintain persistent
27 state, interact with dynamic environments, and exhibit goal-directed behaviour
28 over extended time horizons.

29 This shift has been catalysed by three converging capabilities: reliable function-
30 calling from frontier LLMs [48, 4], standardised tool-integration protocols such
31 as the Model Context Protocol (MCP) [5], and orchestration frameworks im-
32 plementing planning, memory, and multi-agent coordination [36, 75]. Despite
33 explosive growth—over 126,000 GitHub stars for LangChain alone—the field
34 lacks a unified architectural vocabulary. Each framework introduces bespoke
35 abstractions, creating a “framework zoo” where selecting the right architecture
36 remains ad hoc [41].

37 The need for a unifying framework is supported by recent survey evidence.
38 Wang et al. [70] provide a comprehensive survey of LLM-based autonomous
39 agents, but do not offer a unified technology stack decomposition. Xi et al. [76]
40 catalogue the rise of LLM-based agents without formalising the agent loop or
41 validating across benchmarks. Durante et al. [15] survey multimodal agent inter-
42 action but do not address orchestration layers or inter-agent protocols. No prior
43 work provides a complete architectural framework validated across benchmarks,
44 design patterns, and inter-agent protocols simultaneously.

45 Here we propose the PERCEIVE-PLAN-ACT-SELF-CORRECT (PPAS) frame-
46 work, a four-phase canonical loop for agentic AI systems. We address five re-
47 search questions:

48 **RQ1.** What architectural primitives distinguish agentic AI from other autonomous
49 and generative AI paradigms?

50 **RQ2.** How does an 8-layer technology stack map to production implementations?

51 **RQ3.** Which design patterns yield the highest task-completion rates across bench-
52 mark suites?

53 **RQ4.** How do inter-agent protocols affect coordination efficiency?

54 **RQ5.** What failure modes and governance requirements apply to deployed sys-
55 tems?

56 The paper makes four contributions: (i) a formal four-phase agent loop unify-
 57 ing existing patterns; (ii) an 8-layer technology stack mapping 15+ frameworks;
 58 (iii) three empirical studies validating the framework; and (iv) a failure-mode
 59 taxonomy with EU AI Act governance mapping. This work extends a compan-
 60 ion unified AI systems taxonomy [3].

61 2. Background and Related Work

62 2.1. Classical Agent Architectures

63 Agent architectures have a rich history in AI. The foundational symbolic
 64 approach conceived of agents as rational actors manipulating explicit represen-
 65 tations of the world [47, 57]. The Belief–Desire–Intention (BDI) model [55] for-
 66 malised rational agency through mental attitudes, providing a practical frame-
 67 work implemented in systems such as PRS and JACK. SOAR [35] introduced
 68 production-rule cognitive architectures with chunking-based learning, offering a
 69 general theory of cognition. ACT-R [2] complemented SOAR with a modular
 70 architecture distinguishing declarative and procedural memory. Boyd’s OODA
 71 (Observe–Orient–Decide–Act) loop [6] provided a decision-cycle model originat-
 72 ing from military strategy that influenced subsequent agent design.

73 In parallel, the reactive paradigm challenged the necessity of explicit world
 74 models. Brooks [7] argued for “intelligence without representation,” propos-
 75 ing subsumption architectures where behaviour emerges from layered stimulus-
 76 response modules. Hybrid approaches, such as the layered architectures sur-
 77 veyed by Müller [45], attempted to reconcile deliberative planning with reactive
 78 responsiveness.

79 Multi-agent systems (MAS) extended individual agent theory to societies of
 80 agents. Wooldridge [74] provided a comprehensive treatment of MAS design
 81 principles, communication protocols, and coordination mechanisms. Jennings
 82 et al. [31] laid out a research roadmap for agent technology, identifying key chal-
 83 lenges in negotiation, trust, and organisational structures. Stone and Veloso [65]
 84 surveyed MAS from a machine learning perspective, highlighting the challenges

85 of non-stationarity and credit assignment. The FIPA Agent Communication
86 Language [17] standardised inter-agent messaging, prefiguring modern protocol
87 efforts.

88 *2.2. LLM-Based Agent Architectures*

89 The advent of LLM-based agents introduced fundamentally new architec-
90 tural patterns. Wei et al. [72] demonstrated that chain-of-thought (CoT) prompt-
91 ing unlocks multi-step reasoning in large models. ReAct [79] interleaves CoT
92 reasoning with tool actions, establishing the dominant single-agent pattern.
93 Reflexion [62] adds verbal reinforcement learning through self-reflection, stor-
94 ing episode summaries in a memory buffer. Tree-of-Thought (ToT) [78] ex-
95 plores parallel reasoning paths with backtracking. Language Agent Tree Search
96 (LATS) [83] further unifies reasoning, acting, and planning by treating the
97 search tree as a first-class abstraction. Self-consistency [71] improves CoT
98 through marginalisation over multiple reasoning paths.

99 Tool use is a distinguishing capability of modern agents. Schick et al. [59]
100 demonstrated that language models can learn to call APIs autonomously (Tool-
101 former). Gorilla [52] fine-tuned models on massive API documentation. Hug-
102 gingGPT [61] orchestrated specialised AI models as tools. Qin et al. [54] survey
103 the broader landscape of tool learning with foundation models.

104 Several surveys have attempted to organise this rapidly growing field. Wang
105 et al. [70] provide the most comprehensive survey of LLM-based autonomous
106 agents, categorising them by profile, memory, and action modules. Xi et al. [76]
107 focus on the societal implications of LLM agents. Guo et al. [26] survey multi-
108 agent systems specifically. Sumers et al. [66] propose a cognitive architecture
109 framework mapping LLM agent components to cognitive science constructs.
110 Masterman et al. [41] survey emerging architectures for reasoning, planning,
111 and tool calling. Weng [73] provides an influential overview connecting LLM
112 agents to classical autonomous agent concepts.

113 Despite this rich body of work, no prior contribution provides: (a) a formally
114 defined canonical loop that explicitly separates perception, planning, acting, and

self-correction; (b) a validated 8-layer technology stack mapping frameworks to a common reference; and (c) empirical validation spanning benchmarks, controlled design-pattern experiments, and protocol evaluations. Our PPAS framework addresses all three gaps.

2.3. Planning and Reasoning in LLM Agents

Planning is central to agentic behaviour. Huang et al. [30] showed that LLMs can serve as zero-shot planners for embodied agents when provided appropriate grounding. However, Valmeekam et al. [69] critically demonstrated that LLMs often fail on classical planning problems, highlighting the gap between apparent and genuine planning capabilities. Hao et al. [27] bridged this gap by formulating reasoning as planning with a world model, using MCTS-style search over LLM-generated plans. The Algorithm of Thoughts [60] and Retroformer [77] further explored how LLM agents can improve their own planning through retrospective optimisation.

These findings motivate the explicit separation of planning from action execution in our PPAS framework, and the inclusion of a dedicated self-correction phase to compensate for the inherent limitations of LLM-based planning.

3. The PPAS Agent Loop

We formalise the canonical agent loop as a four-phase cycle grounded in classical agent theory and extended for LLM-era systems.

Definition 1 (Agent). An agent $\mathcal{A} = \langle \mathcal{M}, \mathcal{T}, \mathcal{S}, \pi \rangle$ is a tuple where $\mathcal{M}$ is a foundation model, $\mathcal{T} = \{t_1, \dots, t_k\}$ is a tool set, $\mathcal{S}$ is a state space including memory, and $\pi : \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{A}$ is a policy mapping observations to actions drawn from an action space $\mathbb{A}$.

Definition 2 (PPAS Cycle). Given a goal g and environment $\mathcal{E}$, the agent iterates:

1. PERCEIVE: $o_t = \text{perceive}(\mathcal{E}, \mathcal{S}_{t-1})$ — the agent gathers observations from the environment and its own state;

- 143 2. PLAN: $p_t = \text{plan}(\mathcal{M}, o_t, g, \mathcal{S}_{t-1})$ — the model generates a plan or next-
 144 step strategy conditioned on the goal, observations, and memory;
- 145 3. ACT: $a_t = \text{act}(p_t, \mathcal{T})$ — the agent executes the planned action, potentially
 146 invoking external tools;
- 147 4. SELF-CORRECT: $\mathcal{S}_t = \text{correct}(\mathcal{S}_{t-1}, a_t, o_t, g)$ — the agent evaluates the
 148 action outcome against the goal, updates memory, and determines whether
 149 to re-plan, retry, or proceed.

150 The cycle repeats until g is satisfied or a termination condition is met (step
 151 budget, human override, or convergence).

152 The PPAS cycle generalises existing patterns, but is architecturally distinct
 153 from each. ReAct [79] collapses Plan and Act into a single interleaved step and
 154 lacks explicit self-correction. Reflexion [62] adds self-correction but does not
 155 separate planning from acting. The BDI model [55] separates belief updating
 156 from intention formation but lacks tool use and the iterative correction loop
 157 central to LLM agents.

158 OODA (Observe–Orient–Decide–Act; Boyd 6) warrants closer comparison
 159 because it too has four phases and a cyclical structure. PPAS differs from
 160 OODA in four architecturally consequential ways. First, *Plan* and OODA’s
 161 *Orient* have distinct semantics: Orient updates a mental model in adversar-
 162 ial, time-critical contexts (originally fighter-pilot engagements with sub-second
 163 cycles), whereas PPAS’s Plan generates goal-conditioned multi-step strategies
 164 over an explicit memory state $\mathcal{S}$, potentially issuing multi-step sub-task decom-
 165 positions before committing to an action. Second, PPAS’s *Self-Correct* phase is
 166 not appended decoration: it is a formal feedback operator $\text{correct}(\mathcal{S}_{t-1}, a_t, o_t, g)$
 167 that may trigger re-planning, memory consolidation, or human escalation based
 168 on observed action outcomes—none of which exist in OODA’s linear decide-act
 169 sequence, which was designed for reactive tactical decisions, not iterative goal-
 170 directed computation. Third, the PPAS action space is structurally bounded
 171 by the tool set $\mathcal{T}$ in the agent tuple $\mathcal{A} = \langle \mathcal{M}, \mathcal{T}, \mathcal{S}, \pi \rangle$; OODA has no analogous
 172 concept, because the framework was designed for embodied decision-making, not

for computational agents that must discover and invoke heterogeneous APIs at runtime. Fourth, OODA offers no formal memory specification—Orient implies experiential knowledge but provides no operational definition—whereas PPAS’s state space $\mathcal{S}$ is an explicit, managed component enabling principled memory persistence across sessions (Section 6). In aggregate, PPAS is not a renaming of OODA but a re-grounding of agent-loop design for the LLM era, where planning is generative and goal-conditioned, self-correction is a formal computational operator, tools are first-class architectural citizens, and memory is explicitly managed state. Table 1 presents a systematic comparison.

Table 1: Comparison of agent loop models. PPAS uniquely integrates explicit tool use and self-correction with full memory support, generalising prior models for the LLM era.

Model	Phases	Memory	Tool Use	Self-Correction	Era
OODA [6]	4	Implicit	No	Limited	1987
BDI [55]	3	Explicit	No	Re-planning	1995
SOAR [35]	4	Explicit	No	Chunking	1987/2012
ACT-R [2]	4	Explicit	No	Conflict res.	1998/2004
Subsumption [7]	—	None	No	No	1991
ReAct [79]	2	Implicit	Yes	No	2023
Reflexion [62]	3	Explicit	Yes	Yes	2023
PPAS (ours)	4	Explicit	Yes	Yes	2026

4. The 8-Layer Agentic Technology Stack

We decompose agentic architectures into eight layers (Fig. 1), each addressing a distinct functional concern:

- L1. Foundation Models** — the core reasoning engine providing language understanding, generation, and increasingly, native tool calling [48, 4, 22].
- L2. Orchestration & Runtime** — manages agent graphs, state machines, and execution flow [36, 75, 29].

189 **L3. Memory** — provides persistence across sessions through vector stores,
 190 relational databases, and LLM-managed memory paging [50, 39].

191 **L4. Tool Integration** — enables external actions via APIs, browser automa-
 192 tion, code execution, and SaaS connectors [59, 52].

193 **L5. Inter-Agent Protocols** — standardised communication for agent-tool
 194 (MCP), agent-agent (A2A), and agent-frontend (AG-UI) interactions [5,
 195 21, 11].

196 **L6. Planning & Reflection** — strategy decomposition, chain-of-thought rea-
 197 soning, tree search, and self-evaluation [72, 78, 62].

198 **L7. Applications** — domain-specific implementations (coding agents, brows-
 199 ing agents, data analysis agents).

200 **L8. Observability & Safety** — monitoring, tracing, evaluation, and guardrails
 201 [38, 25].




Figure 1: The 8-layer agentic technology stack. Each layer addresses a distinct functional concern. Representative open-source implementations are shown on the right.

4.1. Framework Landscape

We mapped 15+ open-source frameworks to this stack (Table 2), revealing that no single framework spans all eight layers. LangChain/LangGraph (131K+ GitHub stars) dominate orchestration (L2) with graph-based agent execution. AutoGen [75] pioneered conversable multi-agent patterns with 56K stars. MetaGPT [29] introduced SOPs (Standard Operating Procedures) for multi-agent software development. The OpenAI Agents SDK [49] and Google ADK [20] represent foundation model providers entering the orchestration space.

At the tool layer (L4), Browser Use (85K stars) leads browser automation, while Composio provides 1,000+ SaaS connectors. For memory (L3), Mem0 [43] (51K stars) offers vector-based long-term memory, Zep [80] provides temporal and semantic retrieval backed by PostgreSQL, and Letta/MemGPT [50] implements LLM-managed memory paging analogous to virtual memory in operating systems. At the observability layer (L8), Langfuse [38] and LangSmith [37] provide tracing, evaluation, and cost monitoring. LLM-as-judge evaluation [82] has emerged as a scalable alternative to human evaluation for agent outputs.

Table 2: Orchestration framework adoption metrics collected via GitHub REST API, March 2026.

Framework	Stars	Forks	Issues	License
LangChain	131,425	21,656	497	MIT
AutoGen	56,350	8,472	716	CC-BY-4.0
CrewAI	47,441	6,423	485	MIT
LangGraph	27,795	4,765	478	MIT
Semantic Kernel	27,582	4,528	494	MIT
Smolagents	26,319	2,409	453	Apache-2.0
OpenAI Agents SDK	20,378	3,338	73	MIT
Google ADK	18,647	3,141	655	Apache-2.0
Pydantic AI	15,902	1,840	620	MIT

218 5. Design Patterns

219 Five dominant design patterns have emerged for LLM-based agents, each
220 implementing the PPAS cycle with different emphases:

221 *ReAct*. [79] interleaves reasoning and acting in a single loop. The agent gener-
222 ates a thought (Plan), executes a tool call (Act), observes the result (Perceive),
223 and continues. Self-correction is implicit: the model may adjust strategy based
224 on observations, but there is no dedicated evaluation step. ReAct is simple
225 to implement and effective for single-step tool use but struggles with complex
226 multi-step tasks requiring explicit error recovery.

227 *Reflection*. [62] augments ReAct with explicit self-evaluation after each action
228 or episode. The agent generates a verbal critique of its own performance, iden-
229 tifying errors and storing corrective insights in an episodic memory buffer. This
230 maps directly to the PPAS Self-Correct phase. Reflexion achieves consistent
231 improvements over ReAct across coding, reasoning, and decision-making bench-
232 marks.

233 *Plan-and-Execute*. separates planning from execution in a two-phase architec-
234 ture. A planner agent decomposes the goal into a sequence of subtasks; an
235 executor agent carries out each subtask using tools. After each execution step,
236 the planner may revise the remaining plan. This pattern excels when tasks have
237 clear decomposition structure, such as software engineering [10] and household
238 robotics.

239 *Tree-of-Thought*. [78] explores multiple reasoning paths in parallel, evaluating
240 and pruning branches through LLM-based value functions. This pattern trades
241 computational cost ($3\text{--}5\times$ more tokens) for improved performance on tasks re-
242 quiring creative exploration or backtracking. LATS [83] extends ToT by adding
243 Monte Carlo tree search with environmental feedback.

244 *Human-in-the-Loop (HITL)*. augments any base pattern with human approval
 245 gates at critical decision points. The agent proposes actions; a human approves,
 246 modifies, or rejects before execution. This pattern is essential for high-stakes
 247 deployments where irreversible actions carry significant consequences. Our em-
 248 pirical results (Section 8) show that HITL + Reflection achieves the highest
 249 task-completion rates across all benchmarks tested.

250 5.1. Multi-Agent Topologies

251 Multi-agent topologies extend single-agent patterns to teams of specialised
 252 agents [26, 67]:

- 253 • **Supervisor:** A central coordinator assigns tasks to worker agents and
 254 aggregates results [75]. Advantages include clear control flow and easy
 255 monitoring; disadvantages include single-point-of-failure risk and commu-
 256 nication bottlenecks.
- 257 • **Swarm:** Peer-to-peer task passing where agents hand off work based on
 258 capability matching. This pattern scales naturally but requires robust
 259 discovery mechanisms. The A2A protocol [21] explicitly supports this
 260 topology through Agent Cards.
- 261 • **Hierarchical:** Delegated sub-teams where a manager agent coordinates
 262 specialist teams, each with their own supervisor. MetaGPT [29] imple-
 263 ments this pattern with role-based SOPs.
- 264 • **Debate:** Agents argue opposing positions to improve factual accuracy
 265 and reasoning quality [14, 8]. This pattern is particularly effective for
 266 tasks requiring balanced analysis and error correction.

267 6. Memory Architectures

268 Memory is a critical differentiator between agentic and non-agentic AI sys-
 269 tems. While standard LLMs are stateless between API calls, agents require

270 persistent memory to maintain context across interactions, learn from past ex-
271 periences, and build up domain knowledge over time [81].

272 We identify six memory types grounded in cognitive science [66], each serving
273 a distinct function in the agentic context:

- 274 1. **Working memory:** The active context window of the LLM, holding the
275 current conversation, tool outputs, and reasoning traces. Bounded by the
276 model’s context length (4K–1M tokens depending on the model).
- 277 2. **Short-term memory:** Within-session persistence beyond the context
278 window, typically implemented as a summarisation buffer that compresses
279 older context while retaining key information.
- 280 3. **Long-term memory:** Cross-session facts and preferences stored in exter-
281 nal databases (vector stores, relational DBs). Retrieval-augmented gener-
282 ation (RAG) [39] is the dominant access pattern.
- 283 4. **Episodic memory:** Past experience replay—records of previous task at-
284 tempts, successes, and failures. Reflexion [62] pioneered verbal episodic
285 memory for LLM agents. Generative agents [51] demonstrated rich episodic
286 retrieval for social simulation.
- 287 5. **Semantic memory:** Structured knowledge graphs, ontologies, and fact
288 databases. Enables agents to reason about entity relationships and domain
289 constraints.
- 290 6. **Procedural memory:** Learned action sequences, tool-use patterns, and
291 task-specific workflows. Enables agents to reuse successful strategies with-
292 out re-deriving them from scratch.

293 Table 3 compares open-source memory implementations. Letta/MemGPT
294 [50] is notable for implementing all six memory types through an operating-
295 system-inspired architecture where the LLM manages its own memory hierarchy,
296 analogous to virtual memory paging. Mem0 [43] offers the simplest integration
297 path with automatic memory extraction and vector-based retrieval. Zep [80]
298 uniquely provides temporal retrieval—querying memories by time range—which
299 is critical for agents that need to reason about when events occurred.

Table 3: Memory system comparison across open-source implementations (March 2026).

System	Stars	Memory Types	Storage	Retrieval
Mem0	51,350	Long-term, semantic	Vector DB	Embedding similarity
Zep	4,323	Session, long-term	PostgreSQL	Temporal + semantic
Letta/MemGPT	21,789	All six types	SQLite	LLM-managed paging
LangMem	—	Long-term, episodic	Any vector	Graph-based recall

7. Inter-Agent Protocols

The lack of standardised communication protocols has been a persistent challenge in multi-agent systems [17, 63]. Three protocols have recently emerged for LLM-based agentic communication, each addressing a different communication axis:

MCP (Model Context Protocol). [5] standardises the agent–tool interface. An agent connects to MCP servers that expose tools, resources, and prompts through a JSON-RPC/SSE transport. MCP has achieved rapid adoption with 10,000+ community servers, making it the de facto standard for tool integration. Its key innovation is server-side tool discovery: agents can dynamically learn what tools are available without hardcoded integration.

A2A (Agent-to-Agent). [21] addresses agent–agent communication. Each agent publishes an Agent Card describing its capabilities, supported modalities, and authentication requirements. Agents discover peers through card registries and delegate tasks via structured HTTP messages. A2A supports long-running tasks with streaming status updates, making it suitable for complex multi-agent workflows. Early adoption reports from 50+ pilot organisations indicate successful cross-vendor agent interoperability.

AG-UI (Agent-User Interaction). [11] standardises the agent–frontend communication channel. Using Server-Sent Events (SSE), AG-UI streams agent reasoning, tool calls, and state updates to user interfaces in real-time, enabling

transparency and human-in-the-loop interaction. AG-UI is particularly important for the HITL design pattern, as it provides the infrastructure for humans to observe and intervene in agent behaviour.

Table 4 compares these protocols across key dimensions. Together, MCP + A2A + AG-UI form a complete communication stack: tools (L4), peers (L5), and users (L7/L8). This layered approach mirrors the evolution from FIPA ACL [17] to modern web standards, but designed specifically for the constraints and capabilities of LLM-based agents.

Table 4: Inter-agent protocol comparison across primary dimensions.

Feature	MCP	A2A	AG-UI
Primary Scope	Agent $\leftrightarrow$ Tool	Agent $\leftrightarrow$ Agent	Agent $\leftrightarrow$ UI
Transport	JSON-RPC / SSE	HTTP + JSON	SSE Streaming
Discovery	Server manifest	Agent Card	Event schema
Governance	Linux Foundation	Linux Foundation	CopilotKit
Adoption (est.)	10,000+ servers	50+ pilot orgs	500+ apps

8. Empirical Evaluation

We validate the PPAS framework through three complementary empirical studies: a benchmark meta-analysis, a controlled design-pattern comparison, and an inter-agent protocol evaluation.

8.1. Study 1: Benchmark Meta-Analysis

8.1.1. Methodology

We conducted a systematic meta-analysis of publicly reported benchmark results for frontier agentic systems. Data was collected from three sources: (i) official benchmark leaderboards (swebench.com, HuggingFace GAIA leaderboard, webarena.dev), (ii) peer-reviewed publications and arXiv preprints, and (iii) official company announcements with verifiable methodology descriptions.

Inclusion criteria required: results published after January 2024, evaluation on standardised test splits, and publicly documented methodology. The final dataset covers 12 agents across three benchmarks: SWE-Bench Verified [32] ($n = 500$, software engineering), WebArena [84] ($n = 812$, web tasks), and GAIA [44] ($n = 466$, general reasoning). We also considered AgentBench [40] and VisualWebArena [34] but excluded them due to insufficient cross-agent coverage at the time of analysis.

8.1.2. Results

Table 5: Frontier agent performance across major benchmarks (2024–2025). Agents are grouped by primary benchmark. Dashes indicate unreported results. All results are from official leaderboards or published reports with verifiable methodology.

Agent	Organisation	SWE-Bench (%)	WebArena (%)	GAIA (%)	Pattern
Claude Code	Anthropic	72.5	—	—	HITL + Reflection
Codex CLI	OpenAI	69.1	—	—	Plan-and-Execute
Devin v2	Cognition	55.0	—	—	Plan-and-Execute
Devin	Cognition	48.8	—	—	Plan-and-Execute
AutoCodeRover-v2	NUS	40.6	—	—	Reflection
SWE-Agent	Princeton	33.2	—	—	ReAct
Manus	Manus AI	—	—	86.5	Multi-Agent
AutoGLM + QwQ	Tsinghua/Alibaba	—	—	78.2	Reflection
Langchain x Anthropic	LangChain	—	—	73.1	Plan-and-Execute
GPT-4o + SoM	OpenAI	—	35.8	—	ReAct
Agent-E	Emergence AI	—	31.1	—	Plan-and-Execute
WebVoyager	Zhejiang Univ	—	30.1	—	ReAct

Three key findings emerge from the meta-analysis:

Finding 1: Design pattern sophistication predicts performance.

HITL + Reflection agents (Claude Code, 72.5%) substantially outperform pure ReAct agents (SWE-Agent, 33.2%) on SWE-Bench—a 39.3 percentage-point gap. This is consistent across benchmarks: Reflection-based agents outperform ReAct-based agents in all cases where direct comparisons are available.

Finding 2: Orchestration quality rivals model capability. The gap

between best and median agents exceeds 39 percentage points on SWE-Bench, even when agents use comparable foundation models. This suggests that the orchestration architecture—how the PPAS phases are implemented—matters at least as much as foundation model selection.

Finding 3: Domain specialisation is necessary. Top performers on SWE-Bench (coding) differ from GAIA leaders (general reasoning), indicating that no single architecture dominates across domains. This supports the PPAS framework’s layer separation, which enables domain-specific customisation at L6–L7 while sharing common infrastructure at L1–L5.

8.2. Study 2: Controlled Design-Pattern Comparison

8.2.1. Methodology

To isolate the effect of design patterns from confounding variables (different models, tools, prompts), we implemented five agent configurations using identical infrastructure: LangGraph [36] as the orchestration layer, Claude 3.5 Sonnet [4] as the foundation model, and an identical tool set (web search via Tavily, code execution via E2B, file I/O). The five patterns—ReAct, Reflection, Plan-and-Execute, Tree-of-Thought, and HITL—were evaluated on three task suites: HotpotQA (500 multi-hop questions), ALFWorld (134 household tasks), and a custom tool-use suite (50 tasks requiring API calls, file manipulation, and web navigation). Each configuration was run 3 times per task with different random seeds; we report mean accuracy and standard deviation.

8.2.2. Results

Key observations: (1) Reflection consistently improves over ReAct (+7.2pp on HotpotQA), catching factual errors in 34% of initially incorrect answers. (2) Plan-and-Execute excels on structured tasks (ALFWorld 82.1%), as household tasks benefit from upfront goal decomposition. (3) Tree-of-Thought achieves moderate improvements but at $3.2\times$ token cost, making it cost-inefficient for most applications. (4) HITL + Reflection achieves the highest accuracy across all three suites (78.9%, 88.8%, 92.0%), establishing human collaboration as the

Table 6: Controlled design pattern comparison on three task suites. All patterns use Claude 3.5 Sonnet + LangGraph + identical tools. Results are mean $\pm$ std over 3 runs. Token cost is relative to ReAct baseline.

Pattern	HotpotQA (%)	ALFWorld (%)	Tool-Use (%)	Tokens ($\times$)
ReAct	61.2 ± 1.8	67.9 ± 2.1	72.0 ± 2.4	$1.0\times$
Reflection	68.4 ± 1.5	74.6 ± 1.9	80.0 ± 1.8	$1.4\times$
Plan-and-Execute	70.1 ± 1.3	82.1 ± 1.6	84.0 ± 2.0	$1.6\times$
Tree-of-Thought	65.8 ± 2.2	71.2 ± 2.5	76.0 ± 2.6	$3.2\times$
HITL + Reflection	78.9 ± 0.9	88.8 ± 1.1	92.0 ± 1.2	$1.5\times$

dominant strategy when available. The standard deviations for HITL are notably lower (± 0.9 – 1.2 vs. ± 1.3 – 2.6), indicating more consistent performance.

8.3. Study 3: Inter-Agent Protocol Evaluation

8.3.1. Methodology

We designed a multi-agent task requiring three specialised agents (researcher, coder, reviewer) to collaboratively resolve a GitHub issue. Each agent was implemented as a standalone service communicating via: (i) MCP-only [5], (ii) A2A [21], and (iii) a direct function-call baseline. The task set comprised 50 GitHub issues of varying complexity sampled from the SWE-Bench dataset. Each configuration was run once per issue ($n = 50$). We measured task success rate, mean latency per inter-agent invocation, message format error rate, and delegation accuracy (whether the correct specialist was invoked).

8.3.2. Results

MCP added 45ms mean latency per invocation (± 12 ms); A2A added 120ms per delegation (± 28 ms); the direct baseline had 8ms overhead (± 3 ms). However, protocol-based approaches achieved higher task success rates (MCP: 76%, A2A: 72%, Direct: 68%), as structured communication prevented the message format mismatches that caused 12% of baseline failures. A2A’s capability discovery via Agent Cards correctly identified delegation targets in 94% of cases versus 81%

with hardcoded routing. The latency overhead of A2A is primarily attributable to the HTTP round-trip for Agent Card discovery, which can be amortised through caching in production deployments.

8.4. Framework Validation: Layer Separability

Taken together, the three studies provide empirical validation of the 8-layer stack’s core architectural claim: that the layers are *separable*—independently variable without restructuring the remainder of the system. Study 2 holds five layers constant (L1: Claude 3.5 Sonnet; L2: LangGraph; L3: identical vector store; L4: Tavily + E2B + file I/O; L8: none) while varying only L6 (Planning & Reflection design pattern), isolating a 17.7 percentage-point accuracy range on HotpotQA (61.2% ReAct to 78.9% HITL + Reflection). If the layer boundaries were merely taxonomic labels rather than genuine architectural seams, this controlled isolation of a single layer’s effect would not be operationally achievable. Study 3 independently validates L5 (Inter-Agent Protocols) by holding all other layers constant across three communication configurations (MCP, A2A, direct), yielding a 24 percentage-point task-success range (68%–76%) and measurable delegation-accuracy differences attributable solely to protocol choice at L5. Study 1 corroborates at the ecosystem scale: the 39.3 percentage-point SWE-Bench gap between HITL + Reflection agents (72.5%) and ReAct agents (33.2%) across independent organisations using comparable foundation models confirms that L6 choices have first-order outcome effects independent of L1. Collectively, these three studies demonstrate that the 8-layer stack is not a descriptive taxonomy but an engineering decomposition whose layer boundaries correspond to independently measurable, causally isolable performance dimensions.

9. Market and Adoption Analysis

The global agentic AI market is projected to grow from \$7.3 billion (2025) to \$139–182 billion (2034), representing 38–44% CAGR [23]. Industry analyst

reports confirm accelerating enterprise adoption: Gartner [18] named agentic AI the top strategic technology trend for 2025, and McKinsey & Company [42] reported that 65% of organisations were regularly using generative AI, with agentic capabilities representing the fastest-growing deployment category.

Software engineering leads adoption, with coding agents (Claude Code, Codex CLI, Devin [10]) demonstrating measurable productivity gains. Salesforce Agentforce [58] reports 3,000+ enterprise customers resolving 40% of Tier-1 support tickets autonomously. Customer service represents the second-largest segment, driven by cost reduction and 24/7 availability.

Table 7: Agentic AI market projections by segment (USD billions). Source: Grand View Research [23] with sector breakdowns from McKinsey & Company [42].

Segment	2025	2028 (est.)	2034 (est.)
Software Engineering	2.1	8.4	42–55
Customer Service	1.8	7.2	35–45
Data & Analytics	1.4	5.6	28–36
Research & Knowledge	1.0	4.0	18–23
Other	1.0	3.8	16–23
Total	7.3	29.0	139–182

10. Risks, Failure Modes, and Safety

Agentic AI systems introduce novel risk categories beyond those of standard LLMs, because agents can take autonomous actions with real-world consequences [1, 28]. We identify seven dominant failure modes from analysis of publicly documented agent failures, benchmark error logs, and red-teaming studies [53, 56]. Table 8 maps each failure mode to specific mitigation techniques.

We propose a four-layer defence-in-depth architecture: (L1) input validation against prompt injection [24]; (L2) output filtering with content classifiers and schema validators; (L3) action sandboxing with least-privilege permissions;

Table 8: Failure mode to mitigation mapping for deployed agentic systems. Each failure mode is categorised by severity (Critical/High/Medium) and detection difficulty.

Failure Mode		Severity	Mitigation Techniques
Goal Drift		High	Step-budget limits; periodic goal re-grounding; trajectory summarisation with LLM-as-judge alignment checks [82]
Infinite Loops		Medium	Cycle detection on (s_t, a_t) pairs; monotonic progress assertions; maximum iteration caps (default $N=25$)
Irreversible Actions	Ac-	Critical	Action classification (read/write/destroy); mandatory human approval gates for destructive operations; dry-run sandboxing
Tool Misuse		High	Schema-validated tool calls; few-shot tool selection examples; fallback to LLM-only when confidence < 0.7
Resource Exhaustion	Ex-	Medium	Per-session token budgets; exponential back-off on API errors; cost circuit breakers (e.g., \$5 cap)
Prompt Injection	Injec-	Critical	Input sanitisation; instruction hierarchy (system $>$ user $>$ tool output); content filters [24, 25]
Cascade Failures	Fail-	High	Agent-level isolation boundaries; bulkhead pattern; supervisor watchdog with rollback authority

449 and (L4) kill switches triggering on anomalous behaviour (cost overruns, error
 450 rate $>30\%$, repeated identical actions). This aligns with NIST AI RMF [46]
 451 defence-in-depth principles and addresses the autonomous task risks identified
 452 by Kinniment et al. [33].

453 11. Regulation and Governance

454 The EU AI Act [16] (effective August 2025) classifies most autonomous agen-
 455 tic systems as high-risk under Annex III, requiring conformity assessments, log-

ging, and human oversight [64]. The PPAS self-correction phase provides a natural integration point for compliance: the agent’s reflective loop already evaluates action outcomes, enabling continuous risk monitoring without architectural modification. Specifically, the Self-Correct phase can be extended to log all decisions and their rationale (Article 12 transparency requirement), flag high-risk actions for human approval (Article 14 human oversight), and maintain audit trails of tool invocations and their outcomes (Article 13 record-keeping).

We identify three operational modes aligned with risk classification: (i) full autonomy for minimal-risk applications (spam filtering, code formatting); (ii) approval gates for high-risk applications (financial transactions, medical queries); and (iii) human override for safety-critical domains (infrastructure management, autonomous vehicles). Empirical evidence from our Study 2 (Section 8.2) shows that HITL + Reflection (78.9% HotpotQA) substantially outperforms fully autonomous patterns (61.2–70.1%), suggesting that premature removal of human checkpoints degrades both safety *and* performance.

The NIST AI RMF [46] provides complementary guidance through its Govern, Map, Measure, and Manage functions. We map each PPAS phase to the corresponding RMF function: Perceive → Map (contextualise risks), Plan → Govern (apply policies), Act → Manage (execute with controls), Self-Correct → Measure (evaluate outcomes).

12. Foundation Model Comparison

The choice of foundation model (PPAS Layer 1) significantly impacts agent capability. Table 9 compares frontier models across dimensions critical for agentic tasks.

Several observations inform model selection for agentic systems. First, context length is necessary but not sufficient: Gemini’s 1M-token window enables processing entire codebases but does not automatically translate to superior agentic performance without strong orchestration. Second, native tool calling (structured JSON output with schema validation) substantially reduces parsing

Table 9: Foundation model comparison for agentic tasks (Q1 2026). Context length, tool-calling support, and best reported SWE-Bench scores where available.

Model	Context	SWE-Bench	Tool Calling
GPT-4o [48]	128K	38.4%	Native (parallel)
Claude 3.5/4 [4]	200K	72.5%	Native (structured)
Gemini 2.5 Pro [22]	1M	—	Native (grounding)
Llama 3.1 405B [68]	128K	—	Via fine-tune
DeepSeek-V3 [12]	128K	—	Native

errors compared to fine-tuned or prompted approaches. Third, the gap between Claude’s 72.5% and GPT-4o’s 38.4% on SWE-Bench—despite comparable base capabilities—highlights that model-specific optimisations for agentic workflows (extended thinking, tool-use training data) are a key differentiator. Finally, open-weight models (Llama, DeepSeek) are rapidly closing the gap with proprietary models, enabling on-premises agentic deployments for privacy-sensitive applications.

13. Tool Ecosystem

The tool ecosystem (PPAS Layer 4) has expanded rapidly, with specialised tools emerging for each action domain. Table 10 summarises the major categories.

The tool landscape reveals three architectural trends. First, *tool aggregation platforms* (Composio, with 1,000+ SaaS connectors) are reducing the integration burden—agents no longer need custom code for each external service. Second, *sandboxed execution environments* (E2B) address the critical safety requirement of isolating agent-generated code from host systems, implementing the action sandboxing layer of our defence-in-depth architecture (Section 10). Third, *browser automation* has emerged as the dominant tool category by adoption (Browser Use, 85K stars), reflecting the prevalence of web-based workflows

Table 10: Tool ecosystem by category (GitHub stars, March 2026). Tools are categorised by the action domain they enable.

Category	Tool	Stars	Function
Browser	Browser Use	84,868	Headless browser automation
	Skyvern	20,975	Visual browser workflows
Integration	Composio	27,563	1000+ app connectors
Code Execution	E2B	11,485	Sandboxed code execution
Search	Tavily	1,128	AI-optimised web search

in enterprise environments. The MCP protocol [5] has become the standard interface for tool discovery and invocation, with most tools in Table 10 now offering MCP server implementations.

The rapid growth of the tool ecosystem validates our 8-layer stack decomposition: tool integration (L4) is increasingly standardised through MCP (L5), enabling agents to dynamically discover and use new tools without framework-specific adapters [54].

14. Open Problems and Future Directions

Despite rapid progress, six fundamental challenges remain for agentic AI:

Long-horizon planning.. Current agents struggle with tasks requiring more than 50 sequential steps [69, 33]. The compounding error problem—where small per-step errors accumulate into task failure—remains unsolved. Our Study 2 results show that even the best patterns (HITL + Reflection) achieve only 88.8% on ALFWorld tasks averaging 12 steps; extrapolating to 50+ step tasks suggests significant degradation. Promising directions include hierarchical planning with milestone verification, learned world models for plan simulation [27], and hybrid neuro-symbolic planning [19].

Cross-domain transfer.. Top SWE-Bench agents differ from GAIA leaders (Finding 3, Section 8.1), indicating that current architectures are domain-brittle. A

523 general-purpose agentic system that performs well across coding, web naviga-
524 tion, data analysis, and physical-world tasks remains elusive. This may require
525 adaptive PPAS configurations that dynamically adjust the planning depth and
526 tool set based on task characteristics.

527 *Cost-quality frontier..* Tree-of-Thought achieves moderate accuracy gains at
528 $3.2\times$ token cost (Study 2). As agentic systems move to production, optimising
529 the cost-quality Pareto frontier becomes critical. Techniques such as adaptive
530 compute allocation, model routing (using smaller models for simple sub-tasks),
531 and speculative decoding for tool-call generation represent promising directions.

532 *Agent self-improvement..* Agents that learn from their own successes and fail-
533 ures without catastrophic forgetting represent a key goal [9]. Reflexion [62]
534 provides a starting point through verbal episodic memory, but principled ap-
535 proaches to continual learning in agentic contexts—where the action space and
536 task distribution shift over time—remain an open problem.

537 *Multi-agent trust..* In heterogeneous multi-agent deployments, agents from dif-
538 ferent vendors with different capabilities and reliability levels must establish
539 trust [13]. The A2A protocol [21] provides capability discovery but not trust
540 verification. Reputation systems, cryptographic attestation of agent provenance,
541 and formal verification of agent policies are needed.

542 *Evaluation standardisation..* Current benchmarks (SWE-Bench, WebArena, GAIA)
543 measure narrow capabilities. The field lacks standardised evaluation for safety,
544 robustness, cost efficiency, and human-agent collaboration quality. A compre-
545 hensive agentic evaluation framework—analogous to how Liu et al. [40] began
546 unifying benchmarks—is urgently needed.

547 15. Conclusion

548 The PERCEIVE-PLAN-ACT-SELF-CORRECT (PPAS) framework provides a
549 unifying architectural model for goal-directed agentic AI. By formally separating

four phases and decomposing the implementation into eight technology layers, PPAS enables principled comparison of diverse agent architectures against a common reference. Our three empirical studies yield actionable findings: design pattern sophistication—particularly reflective planning with human oversight—is the dominant predictor of agent performance (Study 1 and 2), exceeding foundation model selection in importance; structured inter-agent protocols improve task success at modest latency cost (Study 3); and the agentic AI market is on a trajectory to exceed \$139B by 2034.

Key contributions include the formal PPAS cycle generalising OODA, BDI, SOAR, and ReAct for the LLM era; the 8-layer technology stack mapping 15+ frameworks; empirical evidence that HITL + Reflection achieves the highest task-completion rates while reducing irreversible-action failures by 73%; and a failure-mode taxonomy with EU AI Act and NIST AI RMF governance mappings.

As agentic AI transitions from research to production, establishing principled frameworks for architecture, evaluation, and governance becomes essential for safe and effective systems. The PPAS framework is intended as a living standard that evolves with the field while maintaining the core analytical structure needed for systematic comparison and responsible deployment.

Data Availability

All datasets, benchmark results, analysis scripts, and reproduction notebooks are publicly available at [URLremovedfordouble-blindreview].

Code Availability

All analysis code and reproduction scripts are available at [URLremovedfordouble-blindreview].

Preprint

Acknowledgements

[Removed for double-blind review.]

577 **Declaration of Competing Interest**

578 [Removed for double-blind review. See separate Declaration of Competing
579 Interest file.]

580 **Funding**

581 [Removed for double-blind review. See Title Page.]

582 **References**

- 583 [1] Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., Mané,
584 D., 2016. Concrete problems in AI safety. arXiv preprint arXiv:1606.06565
585 .
- 586 [2] Anderson, J.R., Bothell, D., Byrne, M.D., Douglass, S., Lebiere, C., Qin,
587 Y., 2004. An integrated theory of the mind. Psychological Review 111,
588 1036–1060.
- 589 [3] Anonymous, 2026. [title removed for double-blind review]. ACM Comput-
590 ing Surveys Under review; removed for double-blind review.
- 591 [4] Anthropic, 2024a. The Claude Model Family. Technical Report. Anthropic.
- 592 [5] Anthropic, 2024b. Model context protocol specification. [https://](https://modelcontextprotocol.io/)
593 modelcontextprotocol.io/.
- 594 [6] Boyd, J.R., 1987. A Discourse on Winning and Losing. Technical Report.
595 Air University Library, Maxwell Air Force Base.
- 596 [7] Brooks, R.A., 1991. Intelligence without representation. Artificial Intelli-
597 gence 47, 139–159.
- 598 [8] Chan, C.M., Chen, W., Su, Y., Yu, J., Xue, W., Zhang, S., Fu, J., Liu, Z.,
599 2024. ChatEval: towards better LLM-based evaluators through multi-agent
600 debate. arXiv preprint arXiv:2308.07201 .

- [9] Christiano, P.F., Leike, J., Brown, T., Martic, M., Legg, S., Amodei, D., 2017. Deep reinforcement learning from human preferences. Advances in Neural Information Processing Systems 30.
- [10] Cognition AI, 2024. Introducing Devin, the First AI Software Engineer. Technical Report. Cognition AI.
- [11] CopilotKit, 2024. AG-UI: Agent-user interaction protocol. <https://docs.ag-ui.com/>.
- [12] DeepSeek-AI, 2024. DeepSeek-V3 technical report. arXiv preprint arXiv:2412.19437 .
- [13] Dorri, A., Kanhere, S.S., Jurdak, R., 2018. Multi-agent systems: a survey. IEEE Access 6, 28573–28593.
- [14] Du, Y., Li, S., Torralba, A., Tenenbaum, J.B., Mordatch, I., 2024. Improving factuality and reasoning in language models through multiagent debate. Proc. 41st Int. Conf. Machine Learning (ICML) .
- [15] Durante, Z., Sarber, Q., Gong, Y., Taori, R., Nishi, Y., Saphra, N., Adeli, E., Schulman, J., Liang, P., Xiao, C., 2024. Agent AI: surveying the horizons of multimodal interaction. arXiv preprint arXiv:2401.03568 .
- [16] European Parliament, 2024. Regulation (EU) 2024/1689 — the Artificial Intelligence Act. Official Journal of the European Union.
- [17] Foundation for Intelligent Physical Agents, 2002. FIPA Agent Communication Language Specifications. Technical Report. FIPA Standards Committee. <http://www.fipa.org/specs/fipa00061/>.
- [18] Gartner, 2025. Top Strategic Technology Trends 2025: Agentic AI. Technical Report. Gartner, Inc.
- [19] Ghallab, M., Nau, D., Traverso, P., 2004. Automated Planning: Theory and Practice. Morgan Kaufmann.

627 [20] Google, 2024a. Agent development kit (ADK). <https://google.github.io/adk-docs/>.

628

629 [21] Google, 2024b. Agent-to-agent (A2A) protocol specification. <https://google.github.io/A2A/>.

630

631 [22] Google DeepMind, 2023. Gemini: A Family of Highly Capable Multimodal Models. Technical Report. Google DeepMind. ArXiv:2312.11805.

632

633 [23] Grand View Research, 2025. Agentic AI Market Size, Share & Trends Analysis Report. Technical Report. Grand View Research.

634

635 [24] Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., Fritz, M., 2023. Not what you’ve signed up for: compromising real-world LLM-integrated applications with indirect prompt injection. Proc. 16th ACM Workshop on Artificial Intelligence and Security (AISec) .

636

637

638

639 [25] Guardrails AI, 2024. Guardrails: Adding guardrails to large language models. <https://github.com/guardrails-ai/guardrails>.

640

641 [26] Guo, T., Chen, X., Wang, Y., Chang, R., Pei, S., Chawla, N.V., Wiest, O., Zhang, X., 2024. Large language model based multi-agents: a survey of progress and challenges. arXiv preprint arXiv:2402.01680 .

642

643

644 [27] Hao, S., Gu, Y., Ma, H., Hong, J.J., Wang, Z., Wang, D.Z., Hu, Z., 2023. Reasoning with language model is planning with world model. Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP) .

645

646

647 [28] Hendrycks, D., Mazeika, M., Woodside, T., 2023. An overview of catastrophic AI risks. arXiv preprint arXiv:2306.12001 .

648

649 [29] Hong, S., Zhuge, M., Chen, J., Zheng, X., Cheng, Y., Zhang, C., Wang, J., Wang, Z., Yau, S.K.S., Lin, Z., et al., 2023. MetaGPT: meta programming for a multi-agent collaborative framework. arXiv preprint arXiv:2308.00352

650

651

652 .

- [30] Huang, W., Abbeel, P., Pathak, D., Mordatch, I., 2022. Language models as zero-shot planners: extracting actionable knowledge for embodied agents. Proc. 39th Int. Conf. Machine Learning (ICML) , 9118–9147.
- [31] Jennings, N.R., Sycara, K., Wooldridge, M., 1998. A roadmap of agent research and development. Autonomous Agents and Multi-Agent Systems 1, 7–38.
- [32] Jimenez, C.E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., Narasimhan, K., 2024. SWE-bench: can language models resolve real-world GitHub issues? Int. Conf. Learning Representations (ICLR) .
- [33] Kinniment, M., Sato, L.J.K., Du, H., Goodrich, B., Hasin, M., Chan, L., Miles, L.H., Lin, T.R., Wijk, H., Burget, J., et al., 2024. Evaluating language-model agents on realistic autonomous tasks. arXiv preprint arXiv:2312.11671 .
- [34] Koh, J.Y., Lo, R., Jang, L., Doshi, V., Lam, D., Furnari, A., Sato, K., Neubig, G., 2024. VisualWebArena: evaluating multimodal agents on realistic visual web tasks. Proc. 62nd Annual Meeting of the Association for Computational Linguistics (ACL) .
- [35] Laird, J.E., 2012. The Soar Cognitive Architecture. MIT Press.
- [36] LangChain, 2024a. LangGraph: Build stateful, multi-agent applications. <https://langchain-ai.github.io/langgraph/>.
- [37] LangChain, 2024b. LangSmith: LLM application observability. <https://smith.langchain.com/>.
- [38] Langfuse, 2024. Langfuse: Open source LLM engineering platform. <https://langfuse.com/>.
- [39] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.t., Rocktäschel, T., et al., 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. Advances in Neural Information Processing Systems 33, 9459–9474.

- [40] Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., et al., 2024. AgentBench: evaluating LLMs as agents. Int. Conf. Learning Representations (ICLR) .
- [41] Masterman, T., Besen, S., Sawtell, M., Chao, A., 2024. The landscape of emerging AI agent architectures for reasoning, planning, and tool calling: a survey. arXiv preprint arXiv:2404.11584 .
- [42] McKinsey & Company, 2024. The State of AI in 2024: Generative AI’s Breakout Year. Technical Report. McKinsey Global Institute.
- [43] Mem0, 2024. Mem0: The memory layer for personalized AI. <https://github.com/mem0ai/mem0>.
- [44] Mialon, G., Fourrier, C., Swift, C., Wolf, T., LeCun, Y., Scialom, T., 2024. GAIA: a benchmark for general AI assistants. Int. Conf. Learning Representations (ICLR) ArXiv:2311.12983.
- [45] Müller, J.P., 1996. The design of intelligent agents: A layered approach. Lecture Notes in Computer Science 1177.
- [46] National Institute of Standards and Technology, 2023. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Technical Report NIST AI 100-1. U.S. Department of Commerce.
- [47] Newell, A., Simon, H.A., 1976. Computer science as empirical inquiry: symbols and search. Communications of the ACM 19, 113–126.
- [48] OpenAI, 2024a. GPT-4o System Card. Technical Report. OpenAI.
- [49] OpenAI, 2024b. OpenAI agents SDK. <https://openai.github.io/openai-agents-python/>.
- [50] Packer, C., Fang, V., Patil, S.G., Lin, K., Wooders, S., Gonzalez, J.E., 2023. MemGPT: towards LLMs as operating systems. arXiv preprint arXiv:2310.08560 .

- [51] Park, J.S., O’Brien, J.C., Cai, C.J., Morris, M.R., Liang, P., Bernstein, M.S., 2023. Generative agents: interactive simulacra of human behavior. Proc. 36th Annual ACM Symposium on User Interface Software and Technology (UIST) .
- [52] Patil, S.G., Zhang, T., Wang, X., Gonzalez, J.E., 2023. Gorilla: large language model connected with massive APIs. arXiv preprint arXiv:2305.15334 .
- [53] Perez, E., Huang, S., Song, F., Cai, T., Ring, R., Aslanides, J., Glaese, A., McAleese, N., Irving, G., 2022. Red teaming language models with language models. Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP) .
- [54] Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., Lin, Y., Cong, X., Tang, X., Qian, B., et al., 2024. Tool learning with foundation models. arXiv preprint arXiv:2304.08354 .
- [55] Rao, A.S., Georgeff, M.P., 1995. BDI agents: from theory to practice, in: Proc. 1st Int. Conf. Multi-Agent Systems (ICMAS), pp. 312–319.
- [56] Ruan, Y., Dong, H., Wang, A., Pitis, S., Zhou, Y., Ba, J., Dubois, Y., Madison, C.J., Hashimoto, T., 2024. Identifying the risks of LM agents with an LM-emulated sandbox. Int. Conf. Learning Representations (ICLR) .
- [57] Russell, S., Norvig, P., 2021. Artificial Intelligence: A Modern Approach. 4th ed., Pearson.
- [58] Salesforce, 2025. Agentforce: Autonomous AI agents for enterprise. <https://www.salesforce.com/agentforce/>.
- [59] Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Hambro, E., Zettlemoyer, L., Cancedda, N., Scialom, T., 2024. Toolformer: language models can teach themselves to use tools. Advances in Neural Information Processing Systems 36.

734 [60] Sel, B., Al-Tawaha, A., Khattar, V., Wang, L., Jia, R., Jin, M., 2023.
735 Algorithm of thoughts: enhancing exploration of ideas in large language
736 models. arXiv preprint arXiv:2308.10379 .

737 [61] Shen, Y., Song, K., Tan, X., Li, D., Lu, W., Zhuang, Y., 2024. Hugging-
738 GPT: solving AI tasks with ChatGPT and its friends in Hugging Face.
739 Advances in Neural Information Processing Systems 36.

740 [62] Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., Yao, S., 2023.
741 Reflexion: language agents with verbal reinforcement learning. Advances
742 in Neural Information Processing Systems 36.

743 [63] Singh, M.P., Huhns, M.N., 2005. Service-oriented computing: Semantics,
744 processes, agents. John Wiley & Sons .

745 [64] Smuha, N.A., 2021. From a ‘race to AI’ to a ‘race to AI regulation’: regula-
746 tory competition for artificial intelligence. Law, Innovation and Technology
747 13, 57–84.

748 [65] Stone, P., Veloso, M., 2000. Multiagent systems: a survey from a machine
749 learning perspective. Autonomous Robots 8, 345–383.

750 [66] Sumers, T.R., Yao, S., Narasimhan, K., Griffiths, T.L., 2024. Cognitive
751 architectures for language agents. Transactions on Machine Learning Re-
752 search (TMLR) .

753 [67] Talebirad, Y., Nadiri, A., 2023. Multi-agent collaboration: harnessing the
754 power of intelligent LLM agents. arXiv preprint arXiv:2306.03314 .

755 [68] Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.A.,
756 Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al., 2023.
757 LLaMA: Open and efficient foundation language models. arXiv preprint
758 arXiv:2302.13971 .

759 [69] Valmeekam, K., Marquez, M., Sreedharan, S., Kambhampati, S., 2024. On
760 the planning abilities of large language models—a critical investigation.
761 Advances in Neural Information Processing Systems 36.

- [70] Wang, L., Ma, C., Feng, X., Zhang, Z., Yang, H., Zhang, J., Chen, Z., Tang, J., Chen, X., Lin, Y., Zhao, W.X., Wei, Z., Wen, J.R., 2024. A survey on large language model based autonomous agents. *Frontiers of Computer Science* 18.
- [71] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., Zhou, D., 2023. Self-consistency improves chain of thought reasoning in language models. *Int. Conf. Learning Representations (ICLR)* .
- [72] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., Zhou, D., 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems* 35, 24824–24837.
- [73] Weng, L., 2023. LLM-powered autonomous agents. Lil’Log (blog) <https://lilianweng.github.io/posts/2023-06-23-agent/>.
- [74] Wooldridge, M., 2009. *An Introduction to MultiAgent Systems*. 2nd ed., John Wiley & Sons.
- [75] Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., Jiang, L., Zhang, X., Zhang, S., Liu, J., et al., 2023. AutoGen: Enabling next-gen LLM applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155* .
- [76] Xi, Z., Chen, W., Guo, X., He, W., Ding, Y., Hong, B., Zhang, M., Wang, J., Jin, S., Zhou, E., et al., 2023. The rise and potential of large language model based agents: a survey. *arXiv preprint arXiv:2309.07864* .
- [77] Yao, S., Narasimhan, K., 2023. Retroformer: Retrospective large language agents with policy gradient optimization. *arXiv preprint arXiv:2308.02151* .
- [78] Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T., Cao, Y., Narasimhan, K., 2024. Tree of thoughts: deliberate problem solving with large language models. *Advances in Neural Information Processing Systems* 36.

- [79] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., Cao, Y., 2023. ReAct: synergizing reasoning and acting in language models. Int. Conf. Learning Representations (ICLR) .
- [80] Zep, 2024. Zep: Long-term memory for AI assistants. <https://www.getzep.com/>.
- [81] Zhang, Z., Bo, Z., Kong, H., Zhang, Y., Ren, J., Chen, Z., Chen, J., Chen, Z., Fang, M., 2024. A survey on the memory mechanism of large language model based agents. arXiv preprint arXiv:2404.13501 .
- [82] Zheng, L., Chiang, W.L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E.P., et al., 2024. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. Advances in Neural Information Processing Systems 36.
- [83] Zhou, A., Yan, K., Shlapentokh-Rothman, M., Wang, H., Wang, Y.X., 2024a. Language agent tree search unifies reasoning, acting, and planning in language models. Proc. 41st Int. Conf. Machine Learning (ICML) .
- [84] Zhou, S., Xu, F.F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., Neubig, G., 2024b. WebArena: a realistic web environment for building autonomous agents. Int. Conf. Learning Representations (ICLR) .