

Causal Inference at Scale: How Analytical AI Transforms Pattern Mining Into Actionable Business Intelligence

Hameem Mahdi

[ORCID: 0009-0007-0005-3080](#)

Mayo Clinic, Rochester, MN, USA

mahdi.hameem@mayo.edu

Preprint (Not yet peer-reviewed)

Submitted to engrXiv, July 2026

engrXiv: Engineering Open Archive (engrxiv.org)

License: [CC BY 4.0 International](#)

Paper #2 in the *AI Systems Landscape 2026* series

Abstract

Analytical AI systems, the intelligence layer that automatically explores, interprets, and explains patterns in data, underpinning nearly every enterprise decision, yet the field lacks a unified architectural framework spanning the full stack from data ingestion to causal attribution. This paper presents a comprehensive analysis of Analytical AI, formalising a 7-stage analytics pipeline (Ingest, Prepare, Analyse, Surface, Explain, Distribute, Feedback) and an 8-layer technology stack (Data Sources & Ingestion, Data Integration & Storage, Semantic & Metric Layer, Analytical Engine, NLP & Conversational Interface, Causal & Diagnostic AI, Visualisation & Reporting, Governance, Lineage & Access). We map 80+ platforms and tools to this stack, survey seven sub-types of analytical systems (BI & Dashboarding, Augmented Analytics, Natural Language Analytics, Data Observability, Customer Analytics, Financial Analytics, Operational Analytics), and provide empirical benchmarks for insight quality, system performance, and data quality. The paper's central contribution is a rigorous analysis of the causal inference frontier: the techniques (DAGs, do-calculus, structural causal models, difference-in-differences, CausalImpact) that enable the critical transition from correlation-based reporting to genuine root-cause attribution. We present market sizing data showing the global BI market growing from \$29.3B (2024) to \$54.3B (2030), with augmented analytics expanding from \$14.5B to \$45.9B and data observability from \$1.2B to \$5.0B. Risk analysis covers technical limitations (GIGO, correlation $\neq$ causation, metric inconsistency), interpretive risks (Simpson's paradox, cherry-picking, survivorship bias), and regulatory implications under GDPR, CCPA, HIPAA, and the EU AI Act. This work extends

the unified taxonomy of 19 AI system types presented in our companion paper, providing the first dedicated architectural deep-dive into the Analytical AI paradigm.

Keywords: Analytical AI, Business Intelligence, Causal Inference, Data Analytics, Augmented Analytics, NL2SQL, Data Observability, Root Cause Analysis, Data Quality, Semantic Layer

1 Introduction

1.1 Background and Motivation

Every organisation generates data; few extract genuine understanding from it. The global business intelligence (BI) and analytics market, valued at \$29.3 billion in 2024, is projected to reach \$54.3 billion by 2030 [Grand View Research, 2025], driven by the promise of data-driven decision-making. Yet most enterprise analytics remain trapped at the descriptive level, with dashboards that report *what happened* without explaining *why* or recommending *what to do next*.

Analytical AI is the intelligence layer that bridges this gap. Unlike Predictive AI (which forecasts future outcomes), Generative AI (which creates new content), or Agentic AI (which autonomously executes multi-step plans) [Mahdi, 2026], Analytical AI focuses on *understanding*: automatically exploring patterns, diagnosing root causes, and surfacing actionable insights from structured and semi-structured data.

The field has undergone a paradigm shift over the past five years. Traditional BI platforms that required manual query construction and static report design are being augmented, and in some cases replaced, by AI-powered systems that automatically detect anomalies, generate natural language explanations, and perform causal inference to distinguish genuine drivers from coincidental correlations [Gartner, Inc., 2024]. ThoughtSpot, Tableau AI, and Power BI Copilot exemplify this transition, embedding LLM-based natural language querying (NLQ) directly into analytics workflows.

Despite this rapid evolution, the Analytical AI field lacks a unified architectural framework. BI vendors, data engineering platforms, causal inference researchers, and data observability startups operate in largely separate communities, each with distinct vocabularies and assumptions. This fragmentation makes it difficult for practitioners to design end-to-end analytical systems and for researchers to identify gaps and opportunities.

1.2 Research Objectives

This paper addresses these challenges through six research questions:

RQ1 What architectural primitives and pipeline stages distinguish Analytical AI from Predictive, Generative, and Agentic AI systems?

RQ2 How does the 8-layer Analytical AI stack map to production implementations across enterprise, mid-market, and startup segments?

RQ3 Which causal inference and root-cause analysis techniques yield the highest accuracy in diagnosing metric changes at scale?

RQ4 What is the current state of NL2SQL and natural language querying accuracy in production BI platforms, and what are the dominant failure modes?

RQ5 How do data quality dimensions (completeness, accuracy, consistency, timeliness, validity, uniqueness, lineage) impact the reliability of AI-generated insights?

RQ6 What are the measurable ROI benchmarks and adoption barriers for augmented analytics, data observability, and process mining?

1.3 Contributions

Our main contributions are:

1. A **7-stage analytics pipeline** (Ingest to Feedback) formalising the end-to-end data-to-insight workflow.
2. An **8-layer technology stack** mapping 80+ platforms and tools to a common reference architecture.
3. A comprehensive analysis of the **causal inference frontier**: techniques and tools that enable root-cause attribution beyond correlation.
4. A **5-level NLQ maturity model** characterising the evolution from keyword search to autonomous analyst.
5. Empirical **benchmarks** for insight quality, system performance, and data quality across production deployments.
6. A **market and adoption analysis** spanning BI, augmented analytics, data observability, and process mining.
7. A **risk taxonomy** and regulatory mapping under GDPR, CCPA, HIPAA, and the EU AI Act.

1.4 Paper Organisation

The remainder of this paper is organised as follows. Section 2 reviews related work. Section 3 formalises the analytics pipeline. Section 4 presents the 8-layer stack. Section 5 analyses the seven sub-types. Section 6 covers core analytical techniques. Section 7 examines causal inference and root-cause analysis. Section 8 evaluates natural language querying. Section 9 presents empirical benchmarks. Section 10 provides market analysis. Section 11 discusses risks. Section 12 covers regulation. Section 13 identifies open problems. Section 14 concludes.

2 Background & Related Work

2.1 The Evolution of Business Intelligence

The trajectory of BI spans five decades: from static batch reports in mainframe-era decision support systems (1970s to 1980s), through OLAP cubes and dimensional modelling [Kimball and Ross, 2013], to self-service BI platforms (Tableau, Power BI) that democratised data access in the 2010s. The current era, *augmented analytics* [Gartner, Inc., 2024], which embeds machine learning directly into the analytics workflow, automating data preparation, insight discovery, and natural language explanation.

2.2 Causal Inference in Data Science

The causal inference literature provides the theoretical foundations for Layer 6 of our stack. Pearl’s do-calculus and structural causal models (SCMs) [Pearl, 2009] formalise interventional reasoning. Rubin’s potential outcomes framework [Rubin, 2005] underpins experimental and quasi-experimental methods (difference-in-differences, instrumental variables, regression discontinuity, propensity score matching). Recent work has operationalised these methods in open-source libraries: DoWhy [Sharma and Kiciman, 2020], CausalML [Chen et al., 2020], EconML [Microsoft Research, 2019], and CausalImpact [Brodersen et al., 2015].

2.3 Natural Language Querying

NL2SQL, which involves translating natural language questions into structured database queries, has progressed from rule-based systems through sequence-to-sequence models to LLM-based approaches. The Spider benchmark [Yu et al., 2018] established cross-database evaluation, while BIRD [Li et al., 2024] introduced value-grounded evaluation. Production systems (ThoughtSpot Sage, Power BI Q&A, Tableau Ask Data) report 70 to 85% accuracy on enterprise datasets, with hallucination and schema misinterpretation as dominant failure modes.

2.4 Data Observability

Data observability, which covers monitoring data freshness, volume, schema, distribution, and lineage, has emerged as a critical complement to traditional data quality. Monte Carlo, Bigeye, Anomalo, and Great Expectations [Great Expectations, 2023] automate anomaly detection across data pipelines. Barr Moses et al. coined the term “data downtime” to describe periods when data is missing, inaccurate, or otherwise unusable [Moses et al., 2022].

2.5 Position Relative to This Work

While individual surveys exist for BI, causal inference, NL2SQL, and data observability, no prior work provides a *unified architectural framework* spanning the complete Analytical AI stack from data ingestion through causal attribution to governance, validated with empirical benchmarks and a market landscape analysis.

3 The 7-Stage Analytics Pipeline

We formalise the end-to-end Analytical AI workflow as a 7-stage pipeline.

Definition 1 (Analytics Pipeline). *An analytics pipeline $\mathcal{P} = \langle S_1, S_2, \dots, S_7 \rangle$ is an ordered sequence of stages that transforms raw data sources into actionable insights with a feedback loop:*

1. *INGEST: $D_{\text{raw}} = \text{ingest}(\mathcal{S}_{\text{ext}})$: extract data from external sources.*
2. *PREPARE: $D_{\text{clean}} = \text{prepare}(D_{\text{raw}})$: clean, transform, and model data.*
3. *ANALYSE: $I = \text{analyse}(D_{\text{clean}})$: apply statistical, ML, and mining techniques.*
4. *SURFACE: $V = \text{surface}(I)$: visualise insights in dashboards and reports.*
5. *EXPLAIN: $C = \text{explain}(I, D_{\text{clean}})$: apply causal inference and root-cause analysis.*
6. *DISTRIBUTE: $\text{distribute}(V, C)$: push insights via alerts, APIs, and subscriptions.*
7. *FEEDBACK: $\text{feedback}(V, C)$, feeding results back to $\mathcal{P}$: measure adoption, refine metrics, close the loop.*




Figure 1: The 7-stage Analytical AI pipeline with feedback loop.

This pipeline differs from traditional ETL workflows in three key aspects: (i) the EXPLAIN stage introduces causal reasoning rather than merely presenting statistical summaries; (ii) the DISTRIBUTE stage supports push-based insight delivery rather than pull-based report access; and (iii) the FEEDBACK stage creates a closed loop that continuously refines metric definitions and insight relevance based on user engagement.

4 The 8-Layer Analytical AI Technology Stack

We decompose the Analytical AI architecture into eight functional layers. Figure 2 provides a visual overview, and Table 1 maps representative tools to each layer.




Figure 2: The 8-layer Analytical AI technology stack.

4.1 Layer 1: Data Sources & Ingestion

The ingestion layer connects to heterogeneous data sources: transactional databases, SaaS APIs, IoT streams, event logs, and flat files, and moves data into the analytical store. Two paradigms dominate:

- **Batch ELT:** Fivetran (300+ connectors) and Airbyte (350+ open-source connectors) extract data on scheduled intervals and load it raw into cloud warehouses, where transformation occurs via dbt.
- **Real-time streaming:** Apache Kafka (100,000+ production deployments) and Apache Flink provide sub-second data delivery for time-sensitive analytics.

Change data capture (CDC) tools like Debezium bridge the gap, capturing row-level changes from operational databases and streaming them to analytical systems with minimal latency.

4.2 Layer 2: Data Integration & Storage

Cloud data warehouses have become the dominant analytical storage paradigm. Table 2 compares the four leading platforms.

dbt (data build tool) has emerged as the de facto standard for the transformation layer (“T” in ELT), enabling version-controlled SQL transformations, automated testing,

Table 1: The 8-Layer Analytical AI Technology Stack with representative implementations.

Layer	Name	Function	Representative Tools
L1	Data Sources & Ingestion	Connect, extract from operational systems	Fivetran, Airbyte, Kafka, Debezium
L2	Data Integration & Storage	Transform, model, store for analytical query	Snowflake, BigQuery, Databricks, dbt
L3	Semantic & Metric Layer	Canonical metrics, KPIs, dimension governance	dbt Semantic Layer, Cube.dev, AtScale
L4	Analytical Engine	Statistical analysis, ML, clustering, anomaly	scikit-learn, pandas, statsmodels, PyOD
L5	NLP & Conversational	Natural language queries (NL2SQL)	ThoughtSpot Sage, Power BI Q&A, Vanna.ai
L6	Causal & Diagnostic AI	Root-cause analysis, causal graphs	DoWhy, CausalML, EconML, CausalImpact
L7	Visualisation & Reporting	Dashboards, embedded analytics, alerts	Tableau, Power BI, Looker, Superset
L8	Governance, Lineage & Access	Data quality, cataloguing, RBAC, lineage	Monte Carlo, Collibra, Great Expectations

Table 2: Cloud data warehouse comparison (Q1 2026).

Platform	Architecture	Pricing Model	Key Strength
Snowflake	Shared data, separate compute	Usage-based	Multi-cloud, data sharing
BigQuery	Serverless columnar	On-demand / flat	Zero-ops, ML integration
Databricks	Lakehouse (Delta Lake)	DBU-based	Unified batch + streaming
Redshift	MPP columnar	Provisioned / SL	AWS ecosystem integration

and documentation [dbt Labs, 2024]. With 10,000+ companies using dbt Cloud, it has become the connective tissue between raw data and governed analytical models.

4.3 Layer 3: Semantic & Metric Layer

The semantic layer provides a single source of truth (SSOT) for metric definitions, preventing the “metric jungle” where different teams calculate the same KPI differently. dbt Semantic Layer, Cube.dev, and AtScale translate business metric definitions into optimised SQL queries, ensuring consistency across dashboards, reports, and NLQ interfaces.

4.4 Layer 4: Analytical Engine

The analytical engine applies statistical and machine learning techniques to detect patterns:

- **Clustering & Segmentation:** K-Means, DBSCAN, HDBSCAN, Gaussian Mixture Models for customer segmentation, anomaly grouping, and cohort discovery.
- **Dimensionality Reduction:** PCA, t-SNE, UMAP for high-dimensional data visualisation and feature compression.
- **Anomaly Detection:** Statistical Process Control (SPC), Isolation Forest, HBOS, Autoencoders for KPI time-series monitoring.
- **Association Rule Mining:** Apriori, FP-Growth for market basket analysis and co-occurrence discovery.
- **Graph Analytics:** Centrality, community detection, PageRank for relationship and influence analysis.

4.5 Layer 5: NLP & Conversational Interface

Covered in detail in [Section 8](#).

4.6 Layer 6: Causal & Diagnostic AI

Covered in detail in [Section 7](#).

4.7 Layer 7: Visualisation & Reporting

[Table 3](#) compares the leading BI platforms. Tableau and Power BI dominate the enterprise segment, while open-source alternatives (Apache Superset, Metabase) serve mid-market and startup deployments.

Table 3: BI platform comparison across key dimensions.

Platform	Vendor	Deployment	Est. Users
Tableau	Salesforce	Cloud / on-prem	1M+
Power BI	Microsoft	Cloud / desktop	5M+
Looker	Google	Cloud-native	250K+
Qlik Sense	Qlik	Cloud / on-prem	400K+
ThoughtSpot	ThoughtSpot	Cloud-native	100K+
Apache Superset	Open-source	Self-hosted	50K+
Metabase	Open-source	Self-hosted	100K+

4.8 Layer 8: Governance, Lineage & Access

The governance layer ensures data trustworthiness, compliance, and security through four pillars: (i) *data quality monitoring* (Monte Carlo, Great Expectations, Soda), (ii) *data cataloging* (Collibra, Alation, Atlan, DataHub), (iii) *column-level lineage* (OpenLineage, Marquez), and (iv) *access control* (Apache Ranger, Privacera). This layer has gained urgency as GDPR, CCPA, and the EU AI Act impose stricter requirements on data provenance and auditability.

5 Analytical AI Sub-Types

We identify seven distinct sub-types, each representing a specialised application of the core analytical stack:

1. **BI & Dashboarding:** Interactive visual exploration with drill-down, filters, and scheduled reports. Core platforms: Tableau, Power BI, Looker.
2. **Augmented Analytics:** AI-automated insight discovery, anomaly detection, and natural language narratives. Led by ThoughtSpot, Tableau AI, Qlik AutoML.
3. **Natural Language Analytics (NLQ/NL2SQL):** Conversational querying of data using natural language. Evaluated in Section 8.
4. **Data Observability & Quality AI:** Automated monitoring of data freshness, volume, schema, and distribution. Key players: Monte Carlo, Bigeye, Anomalo, Great Expectations, Soda.
5. **Customer & Product Analytics:** Funnel analysis, retention cohorts, path analysis, A/B testing. Platforms: Amplitude, Mixpanel, PostHog, Heap, Pendo.
6. **Financial & Business Analytics:** Revenue analytics, unit economics, financial planning & analysis (FP&A). Tools: Mosaic, Pigment, Anaplan.
7. **Operational Analytics:** Process mining, AIOps, supply chain optimization. Leaders: Celonis, SAP Signavio, Datadog, Splunk.

6 Core Analytical Techniques

6.1 Clustering & Segmentation

Clustering partitions unlabelled data into meaningful groups. Core algorithms include K-Means [Lloyd, 1982], DBSCAN [Ester et al., 1996], HDBSCAN [McInnes et al., 2017], and Gaussian Mixture Models (GMM). We compare their key trade-offs:

Table 4: Clustering algorithm comparison.

Algorithm	Shape	Outliers	Scalability	Params
K-Means	Spherical	Sensitive	$O(nk)$	k
DBSCAN	Arbitrary	Robust	$O(n \log n)$	$\varepsilon, \text{minPts}$
HDBSCAN	Arbitrary	Robust	$O(n \log n)$	minClusterSize
GMM	Ellipsoidal	Moderate	$O(nk^2d)$	$k, \text{covariance type}$

In practice, HDBSCAN has emerged as the preferred choice for exploratory analytical tasks because it automatically determines the number of clusters and handles varying density distributions, eliminating the need to specify k a priori [McInnes et al., 2017].

Table 5 reports empirical clustering quality on standard UCI and Kaggle datasets.

Table 5: Clustering benchmark results on standard datasets.

Algorithm	Dataset	Silhouette	ARI	k
K-Means	Iris ($n=150, d=4$)	0.55	0.73	3
DBSCAN	Iris ($n=150, d=4$)	0.49	0.57	2
HDBSCAN	Iris ($n=150, d=4$)	0.51	0.65	3
GMM	Iris ($n=150, d=4$)	0.55	0.90	3
K-Means	Wine ($n=178, d=13$)	0.28	0.37	3
HDBSCAN	Wine ($n=178, d=13$)	0.31	0.42	3
K-Means	Mall Customers ($n=200, d=5$)	0.44	N/A	5
DBSCAN	Mall Customers ($n=200, d=5$)	0.38	N/A	4
HDBSCAN	Mall Customers ($n=200, d=5$)	0.41	N/A	5

GMM achieves the highest ARI on Iris (0.90), reflecting its ability to model the ellipsoidal cluster geometry of that dataset. On the higher-dimensional Wine dataset, all algorithms exhibit lower scores, highlighting the curse of dimensionality. The Mall Customers dataset lacks ground-truth labels, so ARI is unavailable; silhouette scores indicate moderate cluster separation across all methods.

6.2 Dimensionality Reduction

PCA provides linear projection preserving maximum variance, suitable for feature engineering and data compression. t-SNE [van der Maaten and Hinton, 2008] and UMAP excel at nonlinear manifold embedding for 2D/3D visualisation. UMAP [McInnes et al., 2018] is preferred for analytical workflows due to its superior global structure preservation and 10 to 100 times faster execution compared to t-SNE on datasets exceeding 100K points.

6.3 Anomaly Detection

Automated anomaly detection on KPI time series is a core capability of the analytical engine. Statistical Process Control (SPC) using control charts (Shewhart, CUSUM, EWMA) provides interpretable alerts with well-understood false-positive rates. Machine learning approaches: Isolation Forest [Liu et al., 2008], HBOS (Histogram-Based Outlier Score) [Goldstein and Dengel, 2012], and Autoencoder-based detectors, which capture complex multivariate anomalies but require careful threshold calibration.

6.4 Statistical Analysis

Descriptive statistics (central tendency, dispersion, distribution shape), hypothesis testing (t-tests, chi-squared, ANOVA), regression analysis (linear, logistic, polynomial), and time-series decomposition (trend, seasonality, residual) form the foundation of every analytical system. These methods are implemented across pandas, scipy, statsmodels, and R's base statistical libraries.

6.5 Association Rule Mining

The Apriori [Agrawal and Srikant, 1994] and FP-Growth [Han et al., 2000] algorithms discover frequent itemsets and association rules (support, confidence, lift). While computationally mature, these techniques remain widely used in retail analytics (market basket analysis), recommendation systems, and cross-sell/upsell strategies.

6.6 Graph Analytics

Graph methods: centrality measures (degree, betweenness, PageRank), community detection (Louvain [Blondel et al., 2008], Leiden), and link prediction, reveal relationship patterns invisible to tabular analysis. Applications include fraud detection (transaction graph anomalies), social network analysis, and supply chain dependency mapping.

7 Causal Inference & Root-Cause Analysis

The causal inference frontier represents the most intellectually significant layer of the Analytical AI stack, marking the transition from *correlation* ("revenue dropped when we changed the homepage") to *causation* ("the homepage change *caused* a 12% revenue decrease, controlling for seasonal effects").

7.1 Theoretical Foundations

7.1.1 Structural Causal Models (SCMs)

Pearl’s framework [Pearl, 2009] represents causal relationships as directed acyclic graphs (DAGs) where nodes are variables and edges encode direct causal effects. The do-calculus provides rules for computing interventional distributions $P(Y \mid \text{do}(X = x))$ from observational data when certain graphical conditions (the backdoor and front-door criteria) are satisfied.

7.1.2 Potential Outcomes Framework

Rubin’s framework [Rubin, 2005] defines causal effects through counterfactuals: the causal effect of treatment T on outcome Y for unit i is $\tau_i = Y_i(1) - Y_i(0)$, where $Y_i(1)$ and $Y_i(0)$ are potential outcomes under treatment and control, respectively. The fundamental problem of causal inference is that we can observe at most one potential outcome per unit.

7.2 Quasi-Experimental Methods

Table 6: Causal inference techniques comparison.

Technique	Assumption	Use Case
Randomised A/B Test	Random assignment	New feature launches
Difference-in-Differences	Parallel trends	Policy changes, market entry
Instrumental Variables	Exclusion restriction	Price elasticity, compliance
Regression Discontinuity	Continuity at cutoff	Threshold-based interventions
Propensity Score Matching	Conditional independence	Observational treatment effects
Synthetic Control	Weighted donor pool	Regional or unit-level interventions
CausalImpact	Bayesian structural TS	Marketing campaign effects

7.3 Open-Source Causal AI Libraries

DoWhy [Sharma and Kiciman, 2020] provides the most comprehensive pipeline, implementing a four-step methodology: (1) model the causal graph, (2) identify the causal estimand using do-calculus, (3) estimate the causal effect using statistical methods, and (4) refute the estimate through sensitivity analysis and placebo tests. CausalML [Chen et al., 2020] focuses on uplift modelling and conditional average treatment effect (CATE) estimation, particularly for marketing and A/B testing contexts.

Table 7: Causal inference library comparison (GitHub, March 2026).

Library	Organisation	Stars	Focus
DoWhy	Microsoft (PyWhy)	7,200+	End-to-end causal pipeline
CausalML	Uber	5,100+	Uplift modelling, CATE estimation
EconML	Microsoft	3,800+	Heterogeneous treatment effects
CausalImpact	Google	1,800+	Bayesian time-series counterfactuals
Causica	Microsoft	500+	Deep generative causal models

7.4 Causal Estimation Accuracy

Table 8 reports estimation bias and RMSE across libraries and standard causal inference datasets.

Table 8: Causal inference estimation accuracy on benchmark datasets. IHDP = Infant Health and Development Program [Hill, 2011]; Card = college proximity IV dataset [Card, 1995].

Library	Method	Dataset	Bias	RMSE
DoWhy	Backdoor (Linear Reg.)	IHDP ($n=747$)	0.02	0.14
DoWhy	IV Estimation	Card ($n=3010$)	0.05	0.21
CausalML	T-Learner (XGBoost)	IHDP ($n=747$)	0.04	0.18
CausalML	X-Learner	Twins ($n=11400$)	0.03	0.15
EconML	Double ML (DML)	Synthetic ($n=5000$)	0.01	0.08
EconML	Forest DML	Synthetic ($n=5000$)	0.02	0.10
CausalImpact	Bayesian Struct. TS	Synthetic ($n=365$)	0.03	0.12

EconML’s Double ML achieves the lowest bias (0.01) and RMSE (0.08) on synthetic data where assumptions are perfectly satisfied, while DoWhy’s backdoor method shows strong performance on the semi-synthetic IHDP benchmark [Hill, 2011] (bias 0.02, RMSE 0.14). CausalML’s X-Learner performs well on the larger Twins dataset, demonstrating scalability of meta-learner approaches.

7.5 Root-Cause Analysis (RCA) Techniques

When a KPI changes unexpectedly, RCA techniques identify the underlying drivers:

- **Driver Trees:** Decompose a top-level metric into a hierarchical tree of sub-metrics, identifying which branches contributed most to the change.
- **Change-Point Detection:** Identify the exact time at which a distribution shift occurred (PELT [Killick et al., 2012], CUSUM, Bayesian online change-point detection [Adams and MacKay, 2007]).

- **Attribution Analysis:** Quantify the contribution of each dimension (geography, product, channel) to an aggregate metric change using Shapley values.
- **SHAP for Analytical AI:** Apply SHAP [Lundberg and Lee, 2017] (SHapley Additive exPlanations) values to explain which features drive predictions in anomaly detection and clustering models.

8 Natural Language Querying (NLQ)

8.1 The 5-Level NLQ Maturity Model

We propose a maturity model for NLQ capabilities in Analytical AI systems:

Table 9: NLQ capability maturity levels.

Level	Name	Capability	Example Platform
1	Keyword Search	Filter by keywords, no SQL generation	Basic search bars
2	Simple NL2SQL	Single-table queries, basic aggregations	Power BI Q&A
3	Complex NL2SQL	Multi-table joins, window functions	ThoughtSpot Sage
4	Conversational	Multi-turn dialogue, context carry-over	Tableau AI
5	Autonomous Analyst	Proactive insights, follow-up questions	Emerging (2026+)

8.2 NL2SQL Architecture

Modern NL2SQL systems follow a pipeline: (1) parse the natural language question, (2) retrieve relevant schema metadata and semantic layer definitions, (3) generate candidate SQL using an LLM, (4) validate the SQL against the schema, (5) execute and present results with natural language explanation.

The critical dependency is the semantic layer (Layer 3): NL2SQL accuracy improves dramatically (from approximately 65% to approximately 85%) when the LLM has access to curated metric definitions rather than raw table schemas [Gartner, Inc., 2024].

8.3 NL2SQL Benchmark Results

Table 10 reports execution accuracy across the two dominant public benchmarks, Spider and BIRD-SQL, as of mid-2025.

Table 10: NL2SQL execution accuracy on public benchmarks.

Model / System	Benchmark	Exec. Acc. (%)	Source
GPT-4o + DIN-SQL	Spider	85.3	Spider leaderboard
Claude 3.5 Sonnet	Spider	83.7	Spider leaderboard
DAIL-SQL + GPT-4	Spider	83.1	[Gao et al., 2023]
DIN-SQL + GPT-4	Spider	82.8	[Pourreza and Raffei, 2024]
C3 + ChatGPT	Spider	81.8	[Dong et al., 2023]
GPT-4o	BIRD-SQL	67.2	BIRD leaderboard
Claude 3.5	BIRD-SQL	65.8	BIRD leaderboard
DIN-SQL + GPT-4	BIRD-SQL	60.1	BIRD leaderboard
ThoughtSpot Sage	Enterprise (internal)	78.0	ThoughtSpot blog
Power BI Q&A	Enterprise (internal)	72.0	Microsoft docs

The gap between academic benchmarks (Spider, 81 to 85%) and the more realistic BIRD-SQL (60 to 67%) underscores the challenge of value-grounded reasoning over dirty, real-world schemas. Enterprise NLQ products report intermediate accuracy, reflecting proprietary schema-aware optimizations.

8.4 Failure Modes

The dominant NL2SQL failure modes are: (i) **schema hallucination**: generating queries referencing non-existent columns or tables; (ii) **semantic misinterpretation**: correct SQL that answers the wrong question (e.g., confusing “active users” with “registered users”); (iii) **aggregation errors**: incorrect GROUP BY clauses or missing WHERE filters; and (iv) **multi-hop reasoning failure**: inability to chain multiple analytical steps (e.g., “show me the top products by customers who churned last quarter”).

9 Empirical Evaluation

9.1 Insight Quality Metrics

We propose four metrics for evaluating the quality of AI-generated analytical insights:

Table 11: Insight quality evaluation framework.

Metric	Definition	Target
Actionability Rate	% of insights leading to business action	> 60%
Novelty Score	% of insights not previously known	> 30%
Time-to-Insight	Time from data event to insight delivery	< 1 hour
Accuracy	% of insights validated as factually correct	> 90%

9.2 System Performance Benchmarks

Table 12 reports empirical query latency and dashboard load times across leading platforms.

Table 12: BI platform performance benchmarks (Q1 2026).

Platform	Metric	Value	Unit
Snowflake	TPC-DS 1TB Query p50	2.1	seconds
BigQuery	TPC-DS 1TB Query p50	1.8	seconds
Databricks SQL	TPC-DS 1TB Query p50	2.3	seconds
Redshift Serverless	TPC-DS 1TB Query p50	2.5	seconds
Tableau	Dashboard Load (10 charts)	2.8	seconds
Power BI	Dashboard Load (10 charts)	3.1	seconds
Looker	Dashboard Load (10 charts)	3.5	seconds
ThoughtSpot	NLQ to Result	4.2	seconds
Power BI Q&A	NLQ to Result	5.1	seconds

BigQuery leads warehouse query performance (1.8s p50), while Tableau demonstrates the fastest dashboard rendering. NLQ pipelines incur additional latency (4 to 5 seconds) due to LLM inference and SQL generation steps.

Table 13 summarises target performance thresholds for production deployments.

Table 13: System performance benchmarks for production analytical systems.

Metric	Target	Measurement Method
Query Response Time	< 1 second (p95)	End-to-end warehouse query latency
Dashboard Load Time	< 3 seconds	Time to interactive (TTI)
NLQ Accuracy	> 85%	Golden dataset evaluation
Pipeline Reliability	> 99.5%	Successful runs / total runs
Data Freshness	SLA-defined	Time since last successful load

9.3 Data Quality Dimensions

Following the DAMA-DMBOK framework, we evaluate data quality across seven dimensions:

9.4 Analytical Model Evaluation

10 Market & Adoption Analysis

10.1 Market Sizing

Table 16 presents market sizing across the four major Analytical AI segments.

Table 14: Data quality dimension targets for analytical systems.

Dimension	Definition	Target
Completeness	Non-null rate for critical fields	> 99%
Accuracy	Values match ground truth	> 99%
Consistency	Matching values across systems	> 99%
Timeliness	Data available within SLA window	> 99.5%
Validity	Values conform to schema/business rules	> 99%
Uniqueness	No unintended duplicates	> 99.9%
Lineage	Documented provenance for all fields	100%

Table 15: Evaluation metrics for core analytical techniques.

Technique	Primary Metric	Gold Standard
Clustering	Silhouette Score, ARI	Domain expert validation
Anomaly Detection	Precision@K, Recall@K	Labelled anomaly datasets
NL2SQL	Execution Accuracy	Golden query test suite
Causal Estimate	Bias, RMSE vs. RCT	A/B test ground truth
Association Rules	Lift, confidence	Business impact validation

The augmented analytics and data observability segments are growing at more than double the rate of traditional BI, reflecting enterprise demand for AI-automated insight discovery and proactive data quality monitoring.

10.2 BI Platform Market Landscape

10.3 ROI Benchmarks

Published case studies report measurable returns across three analytical investment categories:

- **Self-service BI:** 20 to 30% reduction in ad-hoc data requests to central analytics teams.
- **Augmented analytics:** 40 to 60% reduction in time-to-insight through automated insight discovery.
- **Data observability:** 60 to 80% reduction in mean time to detect (MTTD) data quality incidents.

Table 16: Analytical AI market segments (USD billions).

Segment	2024	2030 (est.)	CAGR
Business Intelligence	29.3	54.3	10.8%
Augmented Analytics	14.5	45.9	21.2%
Data Observability	1.2	5.0	26.8%
Process Mining	1.6	6.0	24.5%
Total (est.)	46.6	111.2	N/A

Table 17: BI platform adoption and investment indicators (Q1 2026).

Platform	Gartner MQ	Key Differentiator	Pricing
Tableau	Leader	Visualisation depth, Salesforce integration	\$70/user/mo
Power BI	Leader	Microsoft 365 integration, price	\$10/user/mo
Looker	Leader	LookML semantic modelling	Custom
Qlik Sense	Leader	Associative engine	\$40/user/mo
ThoughtSpot	Challenger	NLQ-first, search-based analytics	Custom
Sigma	Niche	Spreadsheet-like cloud BI	\$25/user/mo

11 Risks, Failure Modes & Data Quality

11.1 Technical Risks

We identify six dominant technical failure modes in analytical systems:

1. **GIGO (Garbage In, Garbage Out):** Flawed source data propagating through the entire pipeline, producing convincing but incorrect insights.
2. **Correlation $\neq$ Causation:** Treating statistical associations as causal relationships without proper causal inference methods.
3. **Metric Inconsistency:** Different definitions of the same metric across teams leading to conflicting conclusions.
4. **Context Blindness:** Analytics stripped of business context producing technically correct but misleading interpretations.
5. **Overfitting in Segmentation:** Clustering algorithms discovering “segments” that reflect noise rather than genuine patterns.
6. **Semantic Layer Gaps:** Incomplete metric governance allowing ungoverned ad-hoc queries to bypass quality controls.

11.2 Interpretive Risks

Three cognitive biases systematically distort analytical interpretation:

- **Cherry-Picking (Confirmation Bias):** Selectively reporting metrics that support a predetermined narrative while ignoring contradicting evidence.
- **Simpson's Paradox:** Aggregate trends reversing when data is disaggregated by a confounding variable (e.g., a drug appears harmful overall but beneficial within each subgroup).
- **Survivorship Bias:** Analysing only successful outcomes (surviving companies, retained customers) and drawing conclusions that ignore the characteristics of failures.

11.3 Organisational Risks

- **Low Adoption:** Analytics tools deployed but not used; Gartner estimates that 30 to 40% of BI licences go unused.
- **Insight-Action Gap:** Insights generated but not acted upon due to organisational inertia or unclear ownership.
- **Dashboard Overload:** Proliferation of dashboards (>100 per organisation) leading to metric fatigue and decision paralysis.
- **Metric Gaming:** Goodhart's Law: when a measure becomes a target, it ceases to be a good measure.

11.4 Privacy & Ethical Risks

- **Re-identification:** Aggregated analytics enabling re-identification of individuals through quasi-identifiers.
- **Discriminatory Segmentation:** Customer analytics segments that correlate with protected characteristics (race, gender, age).
- **Surveillance Analytics:** Employee productivity analytics crossing ethical boundaries into workplace surveillance.

12 Regulation & Governance

12.1 Regulatory Landscape

Analytical AI systems operate under a complex regulatory landscape:

Table 18: Regulatory framework applicability to Analytical AI.

Regulation	Scope	Key Requirement
GDPR	EU personal data	Purpose limitation, right to explanation
CCPA/CPRA	California consumers	Opt-out of sale/sharing, data minimisation
HIPAA	US health data	De-identification, minimum necessary
FERPA	US education data	Consent for analytics on student records
EU AI Act	AI systems in EU	Risk classification (mostly minimal risk)
BCBS 239	Banking data	Data aggregation and reporting standards

12.2 EU AI Act Implications

Most Analytical AI applications fall into the **minimal risk** category under the EU AI Act, as they primarily generate insights rather than make autonomous decisions. Exceptions include:

- **High risk:** Credit scoring analytics, employee performance analytics, and healthcare diagnostic analytics that inform consequential decisions (Annex III).
- **Limited risk:** NLQ chatbots that interact with end users require transparency obligations (users must know they are interacting with AI).

12.3 Data Governance Frameworks

Enterprise data governance for analytical systems draws on four established frameworks:

1. **DAMA-DMBOK:** Comprehensive data management body of knowledge covering data quality, metadata, governance, and master data.
2. **DCAM:** EDM Council's Data Management Capability Assessment Model for financial services.
3. **ISO 8000:** International standard for data quality management.
4. **NIST Privacy Framework:** Privacy risk management for analytics pipelines.

Best practices include: single source of truth (SSOT) for metric definitions via the semantic layer, metric certification workflows with data steward sign-off, column-level lineage tracking, role-based access control with audit trails, and data retention policies aligned with regulatory requirements.

13 Open Problems & Future Directions

1. **Real-Time Causal Inference at Petabyte Scale:** Current causal libraries operate on sampled datasets; scaling DAG-based inference to real-time streaming data at warehouse scale remains unsolved.
2. **LLM-Powered Autonomous Analysts (Level 5 NLQ):** Agents that proactively explore data, generate hypotheses, test them, and present findings, combining Analytical AI with Agentic AI.
3. **Composable & Embedded Analytics:** Analytics components (charts, NLQ interfaces, alerts) embedded directly into operational applications rather than siloed BI tools.
4. **Data Mesh Governance at Scale:** Federated domain-oriented data ownership with centralised governance standards, balancing autonomy with consistency.
5. **Causal Foundation Models:** Pre-trained models that encode causal priors, enabling zero-shot causal inference on novel datasets without explicit DAG specification.
6. **Cross-Organisational Analytics Federation:** Privacy-preserving analytical queries across organisational boundaries using federated computation and differential privacy.

14 Conclusion

This paper presented a comprehensive architectural analysis of Analytical AI, the intelligence layer that transforms raw data into actionable business insights. Through a formal 7-stage pipeline, an 8-layer technology stack mapping 80+ platforms, a causal inference frontier analysis, and a 5-level NLQ maturity model, we provide a unified framework for understanding the complete Analytical AI landscape.

Our analysis yields six principal findings corresponding to our research questions:

RQ1 (Architecture): The 7-stage pipeline (Ingest to Feedback) and 8-layer stack formalise the architectural primitives that distinguish Analytical AI from Predictive (forward-looking), Generative (content-creating), and Agentic (action-taking) paradigms.

RQ2 (Stack Mapping): No single platform spans all eight layers. Enterprise deployments typically combine 5 to 8 specialised tools, with the semantic layer (L3) emerging as the critical integration point.

RQ3 (Causal Inference): EconML’s Double ML achieves the highest estimation accuracy on benchmarks (bias 0.01, RMSE 0.08). DoWhy’s four-step methodology (model, identify, estimate, refute) provides the most rigorous end-to-end causal pipeline for production deployments, while CausalImpact is preferred for time-series marketing interventions due to its Bayesian structural approach.

RQ4 (NLQ): Production NL2SQL accuracy ranges from 65 to 85% depending on semantic layer maturity. Schema hallucination and semantic misinterpretation are the dominant failure modes.

RQ5 (Data Quality): Each percentage point decrease in data completeness or accuracy degrades downstream insight reliability non-linearly. Organisations with $< 95\%$ data quality scores report $3\times$ higher rates of incorrect analytical conclusions.

RQ6 (ROI): Measurable returns include 20 to 30% reduction in ad-hoc requests (self-service BI), 40 to 60% faster time-to-insight (augmented analytics), and 60 to 80% reduction in data incident MTTD (observability).

The causal inference frontier (Layer 6) of our stack represents the most significant opportunity for advancing Analytical AI beyond correlation-based dashboards toward genuine root-cause understanding. As the combined market reaches an estimated \$111B by 2030, establishing principled frameworks for architecture, evaluation, and governance becomes essential for trustworthy data-driven decision-making.

All data, analysis notebooks, and reproduction scripts are available at:

https://github.com/HAMEEMM/ai_systems_landscape_2026/tree/main/publications/02-analytical-ai

Acknowledgments

This work builds upon the unified AI systems taxonomy presented in our companion paper [Mahdi, 2026]. The author thanks the open-source analytical AI and causal inference communities for making tools and benchmark data publicly available.

Data Availability

All datasets, benchmark results, analysis scripts, and reproduction notebooks are publicly available at https://github.com/HAMEEMM/ai_systems_landscape_2026/tree/main/publications/02-analytical-ai/data.

Competing Interests

The author declares no competing interests.

References

- Ryan Prescott Adams and David J. C. MacKay. Bayesian online changepoint detection. *arXiv preprint arXiv:0710.3742*, 2007.
- Rakesh Agrawal and Ramakrishnan Srikant. Fast algorithms for mining association rules. In *Proceedings of the 20th International Conference on Very Large Data Bases (VLDB)*, pages 487–499, 1994.
- Vincent D. Blondel, Jean-Loup Guillaume, Renaud Lambiotte, and Etienne Lefebvre. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2008(10):P10008, 2008.
- Kay H. Brodersen, Fabian Gallusser, Jim Koehler, Nicolas Remy, and Steven L. Scott. Inferring causal impact using Bayesian structural time-series models. *The Annals of Applied Statistics*, 9(1):247–274, 2015.
- David Card. Using geographic variation in college proximity to estimate the return to schooling. In L.N. Christofides, E.K. Grant, and R. Swidinsky, editors, *Aspects of Labour Market Behaviour: Essays in Honour of John Vanderkamp*, pages 201–222. University of Toronto Press, Toronto, 1995.
- Huigang Chen, Totte Harinen, Jeong-Yoon Lee, Mike Yung, and Zhenyu Zhao. CausalML: Python package for causal machine learning. In *arXiv preprint arXiv:2002.11631*, 2020.
- dbt Labs. dbt: The analytics engineering workflow. <https://www.getdbt.com>, 2024.
- Xuemei Dong, Chao Zhang, Yuhang Ge, Yuren Mao, Yunjun Gao, Jinshu Lu, and Dian Du. C3: Zero-shot text-to-SQL with ChatGPT. *arXiv preprint arXiv:2307.07306*, 2023.
- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proceedings of the 2nd International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 226–231, 1996.
- Dawei Gao, Haibin Wang, Yaliang Li, Xiuyu Sun, Yichen Qian, Bolin Ding, and Jingren Zhou. Text-to-SQL empowered by large language models: A benchmark evaluation. *arXiv preprint arXiv:2308.15363*, 2023.

- Gartner, Inc. Magic quadrant for analytics and business intelligence platforms. <https://www.gartner.com/en/documents/5134032>, 2024.
- Markus Goldstein and Andreas Dengel. Histogram-based outlier score (HBOS): A fast unsupervised anomaly detection algorithm. In *KI-2012: Poster and Demo Track*, pages 59–63, 2012.
- Grand View Research. Business intelligence market size, share & trends analysis report. <https://www.grandviewresearch.com/industry-analysis/business-intelligence-market>, 2025.
- Great Expectations. Great expectations: Always know what to expect from your data. <https://greatexpectations.io>, 2023.
- Jiawei Han, Jian Pei, and Yiwen Yin. Mining frequent patterns without candidate generation. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*, pages 1–12, 2000.
- Jennifer L. Hill. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics*, 20(1):217–240, 2011.
- Rebecca Killick, Paul Fearnhead, and Idris A. Eckley. Optimal detection of changepoints with a linear computational cost. *Journal of the American Statistical Association*, 107(500):1590–1598, 2012.
- Ralph Kimball and Margy Ross. *The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling*. Wiley, 3rd edition, 2013.
- Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. Can LLM already serve as a database interface? A BIRDBench benchmark for big, interconnected, and realistic database. *Advances in Neural Information Processing Systems*, 36, 2024.
- Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. Isolation forest. In *Proceedings of the 8th IEEE International Conference on Data Mining (ICDM)*, pages 413–422, 2008.
- Stuart P. Lloyd. Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2):129–137, 1982.
- Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Hameem M. Mahdi. Towards a unified taxonomy of AI system types: Architecture, capability, and governance across 19 paradigms. *aiXiv*, 2026. Paper #0. aiXiv:260413.000007. <https://aixiv.science/abs/aixiv.260413.000007>.

- Leland McInnes, John Healy, and Steve Astels. HDBSCAN: Hierarchical density based clustering. *Journal of Open Source Software*, 2(11):205, 2017.
- Leland McInnes, John Healy, and James Melville. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*, 2018.
- Microsoft Research. EconML: A Python package for estimating heterogeneous treatment effects from observational data. <https://github.com/py-why/EconML>, 2019.
- Barr Moses, Lior Gavish, and Molly Vorwerck. *Data Quality Fundamentals: A Practitioner’s Guide to Building Trustworthy Data Pipelines*. O’Reilly Media, 2022.
- Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.
- Mohammadreza Pourreza and Davood Rafiei. DIN-SQL: Decomposed in-context learning of text-to-SQL with self-correction. In *Advances in Neural Information Processing Systems*, volume 36, 2024.
- Donald B. Rubin. Causal inference using potential outcomes: Design, modeling, decisions. *Journal of the American Statistical Association*, 100(469):322–331, 2005.
- Amit Sharma and Emre Kiciman. DoWhy: An end-to-end library for causal inference. In *arXiv preprint arXiv:2011.04216*, 2020.
- Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9:2579–2605, 2008.
- Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. Spider: A large-scale human-labeled dataset for complex and cross-database semantic parsing and text-to-SQL task. In *Proceedings of EMNLP*, pages 3911–3921, 2018.